

# **GPU algorithms for comparison-based sorting and for merging based on multiway selection**

Diplomarbeit  
von

**Nikolaj Leischner**

am Institut für Theoretische Informatik, Algorithmik II  
der Fakultät für Informatik

Erstgutachter:	Prof. Dr. rer. nat. Peter Sanders
Zweitgutachter:	Prof. Dr. rer. nat. Dorothea Wagner
Betreuender Mitarbeiter:	M.Sc. Vitaly Osipov

Bearbeitungszeit: 2. Februar 2010 – 29. Juli 2010



Ich erkläre hiermit, dass ich die vorliegende Arbeit eigenhändig angefertigt habe und keine anderen als die angegebenen Hilfsmittel verwendet habe.  
(Nikolaj Leischner)



Special thanks go to Sebastian Bauer and Olaf Liebe for helpful discussions, and for enduring my endless rantings about the painful debugging of GPU programs. Further thanks go to my supervisor Vitaly Osipov, foremost for supporting my work and also for giving me motivation when I needed it.

# Contents

<b>1</b>	<b>Abstract (German)</b>	<b>7</b>
<b>2</b>	<b>Introduction</b>	<b>9</b>
2.1	Motivation . . . . .	9
2.2	Thesis structure . . . . .	10
2.3	Work methodology . . . . .	10
<b>3</b>	<b>Basics</b>	<b>11</b>
3.1	GPU architectures . . . . .	11
3.2	GPU programming languages . . . . .	12
3.3	Performance analysis and modeling for GPU algorithms . . . . .	13
3.4	Parallel primitives . . . . .	14
3.5	Programming methodologies for GPUs . . . . .	14
3.5.1	SIMD-conscious vs. SIMD-oblivious algorithm design . . . . .	14
3.5.2	Multiprocessor-conscious vs. -oblivious algorithm design . . . . .	15
<b>4</b>	<b>Related work</b>	<b>15</b>
4.1	Basic strategies for parallel sorting . . . . .	15
4.2	Sorting and merging on GPUs . . . . .	16
4.3	Sorting on multi-core CPUs . . . . .	17
<b>5</b>	<b>GPU merging</b>	<b>17</b>
5.1	Algorithm overview . . . . .	17
5.2	Multiway selection (multi sequence selection) . . . . .	18
5.2.1	Sequential multi sequence selection . . . . .	19
5.2.2	Parallel multi sequence quick select . . . . .	22
5.3	Parallel merging . . . . .	24
5.3.1	Comparator networks . . . . .	24
5.3.2	Odd-even and bitonic merge . . . . .	25
5.3.3	Parallel ranking . . . . .	26
5.3.4	Block-wise merging . . . . .	26
5.3.5	Pipelined merging . . . . .	28
5.3.6	Other parallel merging approaches . . . . .	29
5.4	Sorting by merging . . . . .	30
5.5	Experimental results . . . . .	30
5.5.1	Hardware setup . . . . .	30
5.5.2	Input data . . . . .	30
5.5.3	Multi sequence selection . . . . .	31
5.5.4	Merging networks . . . . .	33
5.5.5	Merging . . . . .	34
5.5.6	Sorting . . . . .	36

<b>6</b>	<b>Revisiting GPU Sample sort</b>	<b>37</b>
6.1	Algorithm overview . . . . .	37
6.2	Algorithm implementation . . . . .	39
6.3	Exploring bottlenecks . . . . .	41
6.3.1	Sorting networks . . . . .	41
6.3.2	Second level sorting . . . . .	42
6.3.3	$k$ -way distribution . . . . .	42
6.4	GPU Sample Sort vs. revised GPU Sample Sort . . . . .	45
6.5	Other possible implementation strategies . . . . .	47
<b>7</b>	<b>Converting from Cuda to OpenCL</b>	<b>47</b>
7.1	The OpenCL programming language . . . . .	48
7.2	An OpenCL implementation of GPU Sample Sort . . . . .	48
<b>8</b>	<b>Conclusion</b>	<b>50</b>
<b>A</b>	<b>Proof of correctness for sequential multi sequence selection</b>	<b>52</b>
A.1	Correctness . . . . .	52
A.2	Consistency and stability . . . . .	53
<b>B</b>	<b>Proof of correctness for block-wise merging</b>	<b>54</b>
<b>C</b>	<b>CREW PRAM Quicksort</b>	<b>55</b>

# 1 Abstract (German)

Sortieren und Mischen sind zwei grundlegende Operationen, die in zahlreichen Algorithmen – deren Laufzeit von ihrer Effizienz abhängt – verwendet werden. Für Datenbanken sind Sortier- und Misch-Algorithmen wichtige Bausteine. Suchmaschinen, Algorithmen aus der Bioinformatik, Echtzeitgrafik und geografische Informationssysteme stellen weitere Anwendungsfelder dar. Zudem ergeben sich bei der Arbeit an Sortier Routinen interessante Muster die sich auf andere algorithmische Probleme übertragen lassen.

Aufgrund der Entwicklungen in der Rechnerarchitektur in Richtung paralleler Prozessoren stellen *parallele* Sortier- und Misch-Algorithmen ein besonders relevantes Problem dar. Es wurden zahlreiche Strategien entwickelt um mit unterschiedlichen Arten von Prozessoren effizient zu sortieren und zu mischen. Trotzdem gibt es nach wie vor Lücken zwischen Theorie und Praxis; beispielsweise ist Coles Sortieralgorithmus trotz optimaler Arbeits- und Schrittkomplexität praktisch langsamer als Algorithmen mit schlechterer asymptotischer Komplexität. Für Sortiernetzwerke die bereits vor 1970 entdeckt wurden hat bis heute niemand einen besseren Ersatz gefunden. Sortier- und Misch-Algorithmen lassen sich auf endlos viele Weisen zu Misch- und Sortier-*Strategien* kombinieren. Aus diesen Gründen bleiben Sortieren und Mischen, sowohl aus praktischer, als auch aus theoretischer Sicht, ein fruchtbares Forschungsgebiet.

Frei programmierbare Grafikprozessoren (GPUs) sind weit verbreitet und bieten ein sehr gutes Verhältnis von Kosten und Energieverbrauch zu Rechenleistung. GPUs beherbergen mehrere SIMD-Prozessoren. Ein SIMD-Prozessor (*single instruction multiple data*) führt eine Instruktion gleichzeitig für mehrere Datenelemente durch. Mit diesem Prinzip lassen sich auf einem Chip mehr Rechenwerke pro Fläche unterbringen als bei einem skalaren Prozessor. Jedoch arbeitet ein SIMD-Prozessor nur dann effizient, wenn es hinreichend viel gleichförmige parallele Arbeit gibt. Zusätzlich verwenden die Prozessoren einer GPU SMT (*simultaneous multi threading*); sie bearbeiten mehrere Instruktionen gleichzeitig um Ausführungs- und Speicherlatenzen zu verdecken. Aus diesen Eigenschaften ergibt sich dass GPUs nur mit massiv parallelen Algorithmen effizient arbeiten können.

In dieser Diplomarbeit beschäftigen wir uns mit GPU-Algorithmen für Sortieren und Mischen. Wir untersuchen eine Parallelisierungsstrategie für die Misch-Operation die bis dato nur für Multi-core CPUs eingesetzt wurde: Selektion in mehreren sortierten Folgen (*multiway selection*, auch bezeichnet als *multi sequence selection*). Für dieses Problem stellen wir einen neuen Algorithmus vor und implementieren diesen sowie einen zweiten, bereits bekannten, Algorithmus für GPUs. Unser neuer Selektionsalgorithmus ist im Kontext einer Misch-Operation auf GPUs praktikabel, obwohl er *nicht* massiv parallel ist. Wir untersuchen verschiedene Techniken um Mischen feingranular zu parallelisieren: Misch-Netzwerke, sowie hybride Algorithmen, welche Misch-Netzwerke in eine sequentielle Misch-Prozedur einbinden. Die hybriden Algorithmen bewahren die asymptotische Komplexität ihrer sequentiellen Gegenstücke und einer von ihnen ist auch der erste Mehrweg-Misch-Algorithmus für GPUs. Wir zeigen empirisch dass Mehrweg-Mischen auf aktuellen GPUs keinen Performance-Vorteil gegenüber Zweiweg-Mischen erzielt – Speicherbandbreite ist beim Sortieren mit GPUs *noch* kein Flaschenhals. Unser bester Misch-Algorithmus, der auf Zweiweg-Mischen basiert, zeigt ein 35% besseres Laufzeitverhalten als vorher veröffentlichte GPU-Implementierungen der Misch-Operation.



Weiterhin setzen wir Optimierungen für unseren zuvor veröffentlichten GPU-Sortierer GPU Sample Sort um und erzielen im Schnitt eine 25% bessere Leistung, in einzelnen Fällen bis zu 55%. Unsere Implementierung verwenden wir zuletzt als Fallstudie um die Portierbarkeit von Cuda-Programmen in die Sprache OpenCL zu evaluieren. Während Cuda lediglich für GPUs von Nvidia verfügbar ist existieren OpenCL-Implementierungen auch für GPUs und CPUs von AMD, IBM CELL, und Intels Knights Ferry (Larrabee). Wir benutzen Nvidias OpenCL-Implementierung um gleichzeitig feststellen zu können ob es praktikabel ist ganz auf Cuda zu verzichten. Aufgrund von instabilen OpenCL-Treibern für aktuelle Nvidia-Hardware führen wir den Vergleich nur auf älteren Nvidia-GPUs durch. Unsere OpenCL-Portierung erzielt dort 88,7% der Leistung unserer Cuda-Implementierung.

Diese Arbeit stellt einen weiteren Schritt in der Erforschung von Sortieralgorithmen und Misch-Algorithmen für datenparallele Rechnerarchitekturen dar. Wir untersuchen alternative Strategien für paralleles Mischen und liefern Bausteine für zukünftige Sortier-Algorithmen und Mehrweg-Misch-Algorithmen. Der in dieser Arbeit vorgestellte Algorithmus für Selektion in mehreren sortierten Folgen ist auch für CPUs geeignet. Unsere Fallstudie zur Portierbarkeit von Cuda nach OpenCL von GPU Sample Sort bietet einen Anhaltspunkt zur Wahl der Implementierungssprache für GPU-Algorithmen.

## 2 Introduction

### 2.1 Motivation

Sorting and merging are two important kernels which are used as subroutines in numerous algorithms, whose performance depends on the efficiency of these primitives. Databases use sort and merge primitives extensively. Computational biology, search engines, real-time rendering and geographical information systems are other fields where sorting and merging large amounts of data is indispensable. Even when their great practical importance is taken aside; many interesting patterns arise in the design and engineering of sorting procedures, which can be applied to other types of algorithms [23].

More specifically, *parallel* sorting is a pervasive topic. Both sorting and merging are problems that lend themselves to parallelization. Numerous strategies have been designed that map sorting successfully to different kinds of processor architectures. Yet there are gaps between theory and practice. Coles parallel merge sort [11], for example, has optimal work complexity and step complexity. In practice, however, it is so far not competitive to simpler but theoretically worse algorithms [30]. *Sorting networks* have been researched since the 1960s [5]. While networks with lower asymptotic complexity are known [3], the 40 years old odd-even network and bitonic network are still in use today because they perform better in practice. Different sorting algorithms can be combined in countless ways into different sorting strategies. This makes parallel sorting a fruitful area for research from both a theoretical and practical point of view.

Changes in processor architecture over the last years further motivate the design of practical parallel algorithms. As the clock rate and performance of sequential processors starts hitting a wall, the trend has become to put more and more processors on a single chip. Multi-core CPUs are commonplace already and a shift to *many-core* CPUs is likely. Memory bandwidth is not expected to increase in the same manner as parallel processing power, effectively making memory access more expensive in comparison to computation and causing processor designers to introduce deeper memory hierarchies. Current generation GPUs already realize these trends. They also provide computational power at a favorable cost and power-efficiency. This makes them an ideal playground for designing and engineering parallel algorithms.

The work on GPU sorting is extensive, see Section 4. However, *merging* on GPUs has to our best knowledge thus far only been considered as a sub-step of sorting routines and as a method for list intersection [13]. In the first part of this thesis we investigate dedicated merging primitives for GPUs.

Previous GPU merging algorithms either use a recursive merging strategy, or divide the work with a distribution step, similar to sample sort. We wondered if multiway selection was a viable strategy for merging on GPUs. Multiway selection (also termed *multi sequence selection*) splits a set of sorted sequences into multiple parts of equal size that can be merged independently. It has previously been employed by multi-core CPU merge sort algorithms. Its main advantage is that it achieves perfect load balancing for the subsequent merge step – an important consideration for algorithms on many-core architectures.

We demonstrate a new multi sequence selection algorithm that is also suitable for multi-core CPUs and give a proof of correctness for it. We implement the algorithm for

GPUs, alongside a different multi sequence selection algorithm. We implement several strategies to expose fine-grained parallelism in merging: bitonic merging, odd-even merging and hybrid strategies that combine sequential two-way and four-way merging with merging networks. The hybrid algorithms maintain asymptotic work efficiency while still making use of data-parallelism. Our merging primitive achieves a higher performance than earlier published merging primitives do. A sort based on our merger proves to be worse than our previously published GPU Sample Sort [25] and worse than a different recent merge sort implementation [35].

That we cannot match the performance of GPU Sample Sort with our merge sort motivates the second part of the thesis: we reconsider our implementation of GPU Sample Sort and improve it incrementally. Our overhauled implementation achieves on average 25% higher performance and up to 55% higher performance for inputs on which our original implementation performs worst.

At the time of writing most researchers choose Cuda for implementing their GPU algorithms. Cuda, however, only supports Nvidia GPUs. Conducting experiments on other hardware platforms like AMD GPUs, Intels upcoming many-core chips or CELL would increase the validity of experimental results. More important, algorithm implementations are of greater practical use when they run on more hardware platforms. OpenCL is supported on AMD GPUs, AMD CPUs, CELL and also available for Nvidia GPUs. We evaluate the current state of Nvidias OpenCL implementation by implementing GPU Sample Sort in OpenCL. Due to a lack of stable OpenCL drivers for recent Nvidia GPUs we are only able to compare an OpenCL implementation to a Cuda implementation on an older Nvidia GPU (G80). With slight code adjustments we are able to reach 88.7% of our Cuda implementations performance.

## 2.2 Thesis structure

This thesis is divided into 8 sections. The first two lay out the motivation for our work and include a short summary. Section 3 introduces the basic concepts of GPU architectures and GPU programming. A discussion of related work follows in Section 4. In Section 5 we discuss the components of our GPU merging routines and provide an extensive experimental analysis. Our previous work [25] is shortly revisited in Section 6. We then explore the bottlenecks of our sample sort algorithm and evaluate an OpenCL implementation of it in Section 7. Lastly, in Section 8, we summarize our results and elaborate on possible future work.

## 2.3 Work methodology

We follow the paradigm of *algorithm engineering*, which is akin to the process by which natural science develops novel theories. Instead of taking a linear path from theoretical design to implementation we *interleave* design and implementation. While complexity theory guides our algorithmic choices, we always implement our algorithms and set up experiments to verify or falsify our choices.

The strong reliance on experiments is motivated by the large gap between theoretical machine models and actual processor architectures. GPUs exhibit properties of both SIMD and SPMD architectures, use bulk-synchronous communication and have a deep

memory hierarchy. While many of the architectural details are vital for performance, it is very difficult to fit all of them into a single theoretical model. Thus, experimentation is mandatory when arguing for an algorithms performance.

## 3 Basics

In this section we describe the most important properties of the processor architectures for which our algorithms are designed. Further, basic concepts and terms for parallel (GPU) algorithms are introduced. In the following, we use the term *GPU* as a synonym for *stream processor* and *many-core processor*.

### 3.1 GPU architectures

GPUs differ significantly from traditional general purpose processors. While the latter are designed to process a sequential stream of arbitrary mixed instructions efficiently, GPUs excel at solving computationally intense data-parallel problems. A GPU consists of several SIMD processing units (also called multiprocessors). Each SIMD unit consists of several ALUs, load/store units, a register file and an additional local memory. Pipelining is used to increase instruction throughput, although in a way that is very different from that used by CPUs; a multiprocessor executes several SIMD groups in parallel (SMT). Therefore it can keep its pipeline filled with instructions that are guaranteed to be independent of each other. This kind of multi-threading has two major implications:

- Branch prediction heuristics, speculative execution, instruction re-ordering and other complex features used in monolithic scalar processors are less important or superfluous.
- Caches are rather employed to save bandwidth than to reduce memory access latencies, since these latencies are already hidden by SMT.

A GPU is connected to a DRAM memory, which, at the time of writing, typically has a size of several GiBs, and which has more than 100 GB/s bandwidth. Access to this memory is expensive: the available bandwidth is shared among tens of processors, and despite multi-threading the high access latencies can only be hidden if there is a sufficient amount of computation per memory access. For the G200-based Nvidia Tesla c1060 the instruction throughput is  $312 \cdot 10^9$  instructions per second, while the memory throughput is  $25.5 \cdot 10^9$  32 bit words per second. Consequently, 12.2 compute instructions per memory instruction are required on this device to rule out bandwidth limitation. For the Fermi-based Gigabyte Geforce GTX480 the instruction throughput is  $676 \cdot 10^9$  instructions per second, and its memory throughput is  $44.3 \cdot 10^9$  32 bit words per second. On this device 15.3 compute instructions per memory instruction are required to rule out bandwidth limitation.

Another property of the GPU off-chip memory is that its transaction size is similar to that of a CPU cache line. This means that memory access with large strides or unordered access patterns can reduce the effective bandwidth to a fraction of the peak bandwidth. The most recent GPUs at the time of writing (Nvidia Fermi and AMD RV870) have two

levels of cache for the off-chip memory, which make partially unordered access patterns less expensive. We refer to the GPUs DRAM as *global memory*.

To give a clearer picture, we enumerate the most important concrete architectural parameters of current generation GPUs:

- 14 - 30 multiprocessors for a combined number of 128 - 1600 ALUs per chip.
- A SIMD width of 32 or 64 and SMT of 2 - 48 SIMD groups per multiprocessor.
- Register files of size  $1024 \times 32$  bit per multiprocessor, up to  $16384 \times 128$  bit per multiprocessor.
- Local shared memory of 16 KiB - 48 KiB per multiprocessor.
- Up to 16 KiB L1 cache per multiprocessor and up to 768 KiB global L2 cache.
- Global memory transaction size of 128 byte.

Analyzing the properties of such complex processor architectures is an interesting topic in itself. We refer to the work of Wong et al. [39] for an in-depth analysis of the Nvidia G200 chip (Geforce GTX280, Tesla c1060), which is our primary development platform for this work – although we carry out most of our final benchmarks on a Geforce GTX480, which is based on the Nvidia Fermi chip.




Figure 1: Nvidia G200 architecture.

## 3.2 GPU programming languages

Currently, two languages are predominant in GPU programming: Nvidias proprietary Cuda [1], and OpenCL [2] from the Khronos Group consortium. Cuda runs exclusively on Nvidia GPUs while OpenCL is also available for AMD GPUs, AMD CPUs and the CELL processor. Both languages introduce identical abstractions at two levels.

**SIMD execution:** The slots of an SIMD instruction are mapped to threads. If one is only concerned by program correctness, then the SIMD property of the hardware can be ignored. As long as there are only small divergences in the execution paths of different threads this is even practical – threads of a SIMD group that must

not execute an instruction are masked out (branch predication). For greater divergences the different paths occurring within an SIMD group have to be serialized. Consequently, processor utilization is reduced by a large degree. Each thread has a private memory which maps to multiprocessor registers. Communication among threads takes place through the local memory of the multiprocessor, which is shared among all threads and which can be synchronized via barrier synchronization.

**Pipelining:** Each thread belongs to a thread block (OpenCL uses the term *work group*). The size of a thread block can and should be greater than the SIMD width of the hardware. If thread blocks are chosen large enough, then the pipelining properties of GPUs with differing SIMD widths can be hidden.

**Processor virtualization:** In addition to the abstractions at the multiprocessor level, the multiprocessors themselves are also virtualized. Depending on register space, shared local memory space and scheduling constraints each multiprocessor is able to execute several thread blocks simultaneously. Further, a GPU program may consist of a much larger number of thread blocks than can be executed concurrently by the GPU. The GPU's runtime environment re-schedules thread blocks until all have finished executing. Therefore a program – given it runs a sufficiently large number of thread blocks – will run efficiently, regardless of the actual number of multiprocessors present on a specific GPU. Thread blocks share access to the DRAM of the GPU. In contrast to the threads of a single block it is impossible to synchronize threads belonging to different thread blocks during the execution of a program. Instead, synchronization takes place implicitly when a GPU program (also called *kernel*) terminates.

Cudas and OpenCLs abstract machine model can be summarized as a set of uniform CREW PRAMs with small private memories, barrier synchronization and a large shared memory to which all PRAMs have access. The machine is programmed in the SPMD (single program multiple data) paradigm. PRAM indices and thread indices are used to break the symmetry of a program.

### 3.3 Performance analysis and modeling for GPU algorithms

In a parallel execution model there are two metrics for measuring the complexity of an algorithm: the number of time steps for a given problem size, and the total amount of work required for a given problem size. At the level of individual SIMD groups it is mostly the step complexity that matters, unless conflicts in memory access lead to serialization. When looking at thread blocks or the whole set of thread blocks of a GPU algorithm we are rather interested in modeling the work complexity. The amount of thread blocks and SIMD groups executed in parallel is typically lower than the total number of thread blocks of a kernel. Conflicting access to both local multiprocessor memory and global DRAM incurs further serialization. Therefore, it would be very difficult to model a realistic step complexity for a GPU algorithm. That said, algorithms with non-optimal work complexity but superior theoretical step complexity may be more efficient in practice; we determine this by experimenting.

In some places we provide instruction count numbers for our algorithms. We count arithmetic instructions, assignments, comparisons, synchronizations and branches. This way of counting is slightly simplified in that we do not differentiate between local memory and registers – variables from local memory may be transparently cached in registers. Our estimates are consequently too optimistic.

### 3.4 Parallel primitives

Our algorithms use two parallel primitives in multiple places: reduction and scan.

**Reduction** Be  $\oplus$  an associative operator that can be evaluated in constant time. Compute  $\bigoplus_{i < n} x_i := (\dots((x_0 \oplus x_1) \oplus x_2) \oplus \dots \oplus x_{n-1})$ . In the PRAM model  $n$  processors can perform a reduction of  $n$  elements in  $\mathcal{O}(\log(n))$  steps with  $\mathcal{O}(n)$  work. A parallel reduction is therefore work-efficient, but step-inefficient (the number of active processors halves with each step). Step-efficiency can be obtained by pipelining multiple reductions, or by choosing the number of processors  $p \ll n$  and letting each processor perform a part of the reduction serially.

**Scan** Be  $\oplus$  an associative operator that can be evaluated in constant time. An inclusive scan is defined as  $\text{scan}_{\text{in}}(x_0 \dots x_n, \oplus) := \bigoplus_{i \leq 0} x_i, \bigoplus_{i \leq 1} x_i, \dots, \bigoplus_{i \leq n} x_i$ . An exclusive scan is defined as  $\text{scan}_{\text{ex}}(x_0 \dots x_n, \oplus) := \text{id}, \bigoplus_{i < 1} x_i, \dots, \bigoplus_{i < n} x_i$ , where  $\text{id}$  is the identity element of the operator  $\oplus$ .

Blelloch gave a systematic discussion of the scan (also termed *prefix sum*) primitive [6]. Its applications range from lexical analysis to solving recurrences [7]. In the context of parallel sorting scans are commonly used to determine the global address to which an element must be written after some local computation.

Various parallel algorithms for scan exist, with different trade-offs between step complexity and work complexity. Algorithms with optimal work complexity take  $\mathcal{O}(n)$  work for a scan of  $n$  elements, step-optimal algorithms take  $\mathcal{O}(\log(n))$  steps to scan  $n$  elements with  $n$  processors (but require  $\mathcal{O}(n \log(n))$ ). A thorough discussion of scan implementations for GPU architectures is given by Merrill and Grimshaw [27]. Where possible we use existing implementations of scan primitives instead of building our own.

### 3.5 Programming methodologies for GPUs

#### 3.5.1 SIMD-conscious vs. SIMD-oblivious algorithm design

It is possible to ignore the SIMD properties of GPUs and treat a thread block like a PRAM. This helps in designing algorithms that are simple and easy to port to other kinds of parallel architectures. We call such algorithms SIMD-oblivious. It can be advantageous, however, to make conscious use of SIMD execution. If communication between threads is constrained to threads that belong to the same SIMD group, then it is possible to omit explicit synchronization as the threads of a group execute synchronously in any case. Where frequent synchronization is needed this improves processor utilization: whenever a synchronization barrier is hit, fewer and fewer SIMD groups will be able to execute until all groups have reached it – the multiprocessors pipeline will run

empty. The downside to SIMD-conscious programming is that algorithms become more complicated and hardware-specific. When working without explicit synchronization subtle errors can be introduced. E.g. when diverging execution paths within an SIMD group are serialized and the implementations behavior unintentionally depends on the order in which these paths are executed. SIMD-conscious programming is unnecessary if synchronization barriers are infrequent and threads have a uniform workload. Oblivious algorithms are more desirable since they tend to be less hardware-specific and easier to implement. Nevertheless, to use GPUs efficiently, being conscious of SIMD execution is at times unavoidable.

### 3.5.2 Multiprocessor-conscious vs. -oblivious algorithm design

Because processors are virtualized, the actual number of multiprocessors present on a GPU can be ignored in algorithm design and implementation. When parallelization introduces a significant overhead one strives to split a task among as few thread blocks as possible. It is then more efficient to launch just enough thread blocks to saturate a given GPU than to use very many thread blocks. The downside to this approach is that it might lead to hardware-specific design choices; a strategy might work efficiently on the GPUs it was developed and tested on but break down on other architectures. For example when significantly more thread blocks are required to saturate these architectures. Another practical concern is that it is difficult to determine optimal launch parameters at run-time without making queries for device names or similar. The Cuda API provides some functionality for this – multiprocessor count, register file size and other parameters are known at run-time. Although the device information available in Cuda is more extensive than it is for OpenCL it is still insufficient in some cases. The number of registers used by a GPU program influences how many thread blocks can be run at the same time. Information about register usage is not available at run-time and may change between different compiler versions. In summary, multiprocessor-conscious programming leads to better hardware utilization and lower parallelization overheads but makes implementations more hardware-specific and difficult to maintain.

## 4 Related work

Since sorting is one of the most widely studied areas in computer science, there is too much work done in the area to review it exhaustively. Thus, we focus on work that is directly relevant to ours, i.e. sorting and merging on many-core architectures. In the following the terms *GPU* and *many-core architecture* are used interchangeably.

### 4.1 Basic strategies for parallel sorting

Two approaches are commonly used for comparison-based parallel sorting.

**Merging.** A  $k$ -way merging algorithm first sorts equal-sized tiles of the input locally. Subsequently, groups of  $k$  sorted tiles are merged recursively, forming larger sorted tiles, until in the last step the whole input is sorted.



**Distribution.** The other approach –  $k$ -way distribution – proceeds in the opposite direction. Keys are ordered globally: they are distributed to  $k$  buckets. Each key assigned to bucket  $b_i$  is larger than any key in the buckets  $b_0 \dots b_{i-1}$  and smaller than any key in the buckets  $b_{i+1} \dots b_k$ . This procedure is applied recursively to the buckets. Concatenating the sorted buckets yields the result. The sizes of the buckets resulting from a distribution step depend on the distribution of the input data. If buckets are to be sorted in parallel load balancing becomes an issue, as their size may vary greatly.

Quicksort and two-way merge sort are the basic forms of comparison-based distribution and merge sorting. A major reason for constructing more sophisticated algorithms that use a larger value for  $k$  is that this results in fewer scans of the input, reducing the number of memory accesses by a constant factor. Radix sort is a popular distribution sort strategy that is based on the bit-wise representation of keys instead of comparisons.

Most comparison-based algorithms switch to different strategies when a tile or bucket completely fits into a processors cache, local memory and registers. For this task sorting networks [5], e.g. bitonic sort and odd-even merge sort, are a popular choice. While their asymptotic work complexity of  $\mathcal{O}(n \log n \log n)$  is not optimal, they are very space-efficient and inherently parallel.

## 4.2 Sorting and merging on GPUs

Haris et al. were the first to explore sorting on generally programmable GPUs. They developed a radix sort and a hybrid radix-merge sort, making use of the prefix sum (also called *scan*) primitive [36] – refer to Section 3.4 for a description of scan. Based on a segmented scan they also created a parallel quicksort, albeit with impractical performance. Hybridsort by Sintorn and Assarsson [37] uses a single distribution step which assumes uniformly distributed input data – and requires numerical keys. Then it applies a two-way merge procedure to each resulting bucket in parallel, where one thread always merges two tiles. Thus, the last merging steps only exhibit little parallelism if the bucket sizes vary too much. Cederman and Tsigas [9] created a competitive quicksort using a three level strategy: at first a global quicksort distributes the input data to a sufficient number of buckets. Each bucket is then sorted on a multiprocessor with a local quicksort, which applies bitonic sort for buckets fitting into a multiprocessors local memory. Ding et al. implemented a merging primitive based on Coles merge sort [13]. As part of the Thrust and CUDPP libraries Haris et al. [34] published another radix sort algorithm and a two-way merge sort based on the concept by Hagerup and Rüb [18]. Another approach is Bbsort. Similar to Hybridsort, this algorithm assigns numerical keys to buckets, assuming uniformly distributed input. Bitonic sort is applied to buckets that fit into a multiprocessors local memory. Both Peters et al. [31] and Baraglia et al. [4] have designed implementations of bitonic sort which reduce access to global GPU memory by a constant factor. Dehne and Zaboli published a deterministic sample sort [12], which achieves the same performance as our randomized sample sort [25]. Another parallel merge sort, which combines two-way merging and merging networks, was designed by Ye et al. [40]. Ha et al. [17] and Huang et al. [20] developed variants of radix sort which outperform the CUDPP radix sort implementation. Merril and Grimshaw [27, 28]

continued the work on scan primitives for GPU architectures and used these to implement radix sort. When considering 32 bit key-length inputs Merrill and Grimshaws radix sort outperforms all other GPU sorting implementations known to us. Satish et al. [35] presented a GPU merge sort, based on a selection algorithm by Francis et al. [15], which outperforms all other comparison-based GPU sorting algorithms, and a radix sort which delivers – on Intels Knights Ferry (Larrabee) – performance similar to what Merrill and Grimshaw achieve on the Geforce GTX280.

### 4.3 Sorting on multi-core CPUs

To put these GPU algorithms into perspective we hint at recent developments regarding sorting on multi-core CPUs. A sorting routine which is part of the GCC multi-core STL [32] divides the input into  $k$  tiles, where  $k$  is equal to the number of processors. Each processor locally applies introsort [29] to its tile. Finally, the intermediate results are cooperatively  $k$ -way merged. Another multi-core algorithm [10] also follows the pattern of  $k$ -way merging. However, it exploits SIMD instructions and uses bitonic sort [5] for tiles that fit into processor cache. The fastest published multi-core CPU sort is a radix sort by Satish et al. [35].

## 5 GPU merging

In this section we describe the design of a merging primitive for GPUs and give an empirical evaluation of its performance.

### 5.1 Algorithm overview

We intend to merge  $k \geq 2$  sorted sequences with a combined length of  $n$  into one sorted sequence. The original sequences can be of different length. As motivated in Section 3.2, we must decompose the problem of merging at two levels: several thread blocks need to work on it in parallel and within each thread block the work has to be split among multiple threads.

It is clear that the naive approach to parallel merging – merging all pairs of sequences in parallel – does not suffice. We would have to assume unrealistic large values for  $k$  to achieve enough parallelism during the first steps of merging. Great variances in sequence length would cause poor load balancing. Instead, we adopt a strategy that is used by CPU merging algorithms [32]. Prior to merging, a multiway selection (also termed *multi sequence selection*) routine determines  $p - 1$  splitters for each sequence such that the position of the globally smallest  $\frac{n}{p}, \frac{2 \cdot n}{p} \dots \frac{(p-1) \cdot n}{p}$  elements is known. The  $k \times p$  parts then can be merged independently by  $p$  thread blocks. Each thread block gets the same number of elements to merge in total. Generating enough splitters to allow each thread to merge in parallel would be inefficient. We use other techniques to allow the threads within the individual thread blocks to merge sequences in parallel.

Figure 2 shows a high level view of the merging primitive. In the following we explain the algorithms for multi sequence selection and parallel merging.

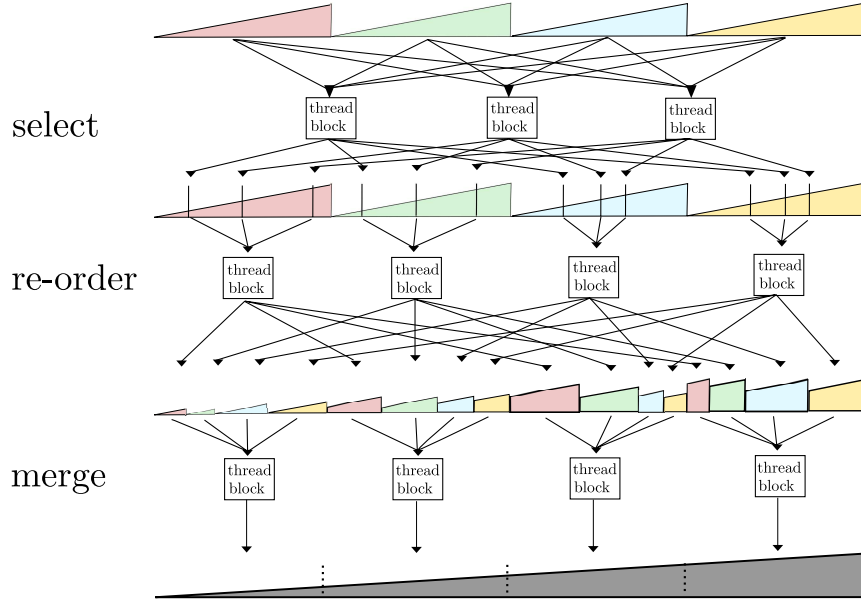


Figure 2: Merging four sorted sequences with our algorithm.

## 5.2 Multiway selection (multi sequence selection)

Given  $k$  sorted sequences of length  $m$  and of combined length  $k \cdot m := n$ , find splitters  $\text{split}_1 \dots \text{split}_m$  such that the sequences are partitioned into the  $f \cdot n$  smallest elements and the  $(f - 1) \cdot n$  largest elements,  $f \in [0, 1]$ . Varman et al. gave an algorithm with a time complexity of  $\mathcal{O}(k \log(fm))$  for this problem [38]. Due to its large constant factors it was never implemented, at least to our knowledge. In the same paper Varman et al. describe a simpler variant with a complexity of  $\mathcal{O}(k \log(k) \log(fm))$ . This second algorithm is used by CPU merging routines.

Because we use multi sequence selection to partition sequences into multiple parts, the selections must be *consistent*: a selection has to include all elements that any smaller selection would include. This is no issue if the keys in the input sequences are unique; selections are then unique as well. But if the greatest key that is included in a selection is not unique and occurs in more than one input sequence, then several combinations of splitters exist which form a valid selection. Figure 3 illustrates selection consistency.




Figure 3:  $\{(1, 0), (2, 0)\}$  – the first selection includes 1 element from the first sequence and none from the second, the second selection includes 2 elements from the first sequence and none from the second – is a consistent selection-set. We call it consistent because every selection in the set includes all smaller selections. This set partitions the sequences in 2 parts of size 1 and 1 part of size 4. In contrast,  $\{(1, 0), (0, 2)\}$  is inconsistent.

Another concern in the presence of duplicate keys is the stability of merging. We must select duplicate keys in the order in which they appear in the input, otherwise their relative order may be changed by our partitioning.

We implement two algorithms. The first is due to Vitaly Osipov and akin to the simpler of the two selection routines devised by Varman et al. The second one lends itself better to parallelization but has an expected work complexity of  $\mathcal{O}(k \log(m) \log(m))$  – for uniformly distributed inputs.

Alternative approaches to multi sequence selection include the recursive merging of splitters used by Thrusts merge sort [34], and random sampling [40]. The latter method may create partitions of unequal size and therefore may cause load imbalance.

### 5.2.1 Sequential multi sequence selection

In Algorithm 1 we maintain a lower boundary, upper boundary and a sample element for each sequence. The sample position and lower boundary are initialized to 0, while the upper boundary is initialized to  $m$ . In addition, a global lower and upper boundary is stored, indicating how many elements are already selected at least and how many are selected at most. If the global upper boundary is smaller than the targeted selection size, then the sequence with the smallest sample element is chosen. If the sample element of more than one sequence has the smallest value, then the sequence with the lowest index is chosen. The selection from this sequence is increased, and all boundaries are adjusted accordingly. Otherwise, the sequence with the greatest sample element is selected and the selection from it is decreased until the greatest selected element in it is smaller than the greatest element of which it is known that it belongs to the selection. If the sample element of more than one sequence has the greatest value, then the sequence with the highest index is chosen. This process is repeated until the global boundaries match. Figure 4 illustrates the process. Breaking ties by sequence index ensures consistent selections that are also suitable for stable merging. A proof of correctness for Algorithm 1 is provided in Appendix A.




Figure 4: Algorithm 1 gradually refines its selection in both sequences.

The sample selection is done with a priority queue, which can be implemented e.g. as a tournament tree. In our GPU implementation we use a parallel reduction instead – refer to Section 3.4 for a description of the reduction primitive. While a reduction has a work complexity of  $\mathcal{O}(n)$  instead of the  $\mathcal{O}(\log(n))$  achieved by a tournament tree, the

latter is completely sequential and therefore a bad match for GPU architectures. Also, in practice we use parameters for which the reduction is done by a single SIMD group. As established in Section 3.3 it is then step complexity which matters. For the rest of the algorithm we have no choice but to execute it sequentially. In Section 5.5 we show that the sequential execution is affordable when using multi way selection as a sub-step of merging.

---

**Algorithm 1:** Sequential multi sequence selection.

---

```
/* Global boundaries and sequence boundaries */
 $U := L := 0$   $l_1 \dots l_k := 0$   $u_1 \dots u_k := m$ 
/* Samples from all sequences and sample positions. */
 $s_1 := \text{seq}_1[0] \dots s_k := \text{seq}_k[0]$   $p_1 \dots p_k := 0$ 
repeat
  if  $U \leq t$  then
     $i := \text{getMinSampleSequenceId}(s)$ 
     $x := s_i$ 
    /* Binary search step, increase selection size. */
     $n := \frac{u_i + p_i + 1}{2}$ 
    if  $p_i \neq 0$  then
       $L := L + p_i - l_i$ 
       $U := U + n - p_i$ 
    else
       $L := L + 1$ 
       $U := U + n + 1$ 
     $s_i := \text{seq}_i[n]$ 
     $l_i := p_i$ 
     $p_i := n$ 
  else
     $i := \text{getMaxSampleSequenceId}(s)$ 
     $o := l_i$ 
     $\text{low} := \frac{l_i + p_i}{2}$   $\text{high} := p_i$ 
    /* Binary search, decrease selection size. */
    while  $\frac{\text{low} + \text{high}}{2} \neq \text{low}$  and  $\text{seq}_i[\text{low}] < x$  do
       $l_i := \text{low}$ 
       $\text{low} := \frac{\text{low} + \text{high}}{2}$ 
    end
     $L := L + l_i - o$ 
     $U := U + \text{low} - p_i$ 
    if  $\text{low} \neq o$  then
       $u_i := \text{high}$ 
       $s_i := \text{seq}_i[\text{low}]$ 
       $p_i := \text{low}$ 
    else
       $u_i := \text{low}$ 
until  $U = L$ 
```

---

### 5.2.2 Parallel multi sequence quick select

Due to its sequential nature Algorithm 1 is not particularly GPU-friendly. Therefore we investigate an alternative approach, which exposes more parallelism. Its downside is suboptimal work complexity.

As in Algorithm 1, in Algorithm 2 we maintain a lower and upper boundary for each sequence. In each step we pick a pivot element uniformly at random from one sequence. The sequence from which we pick is also selected uniformly at random. We rank the pivot in each sequence via binary search. If the sum of all ranks is larger than the desired selection size we use the ranks as new upper boundaries, otherwise we adjust all lower boundaries. This process is repeated until the correct number of elements is selected.

---


**Algorithm 2:** Parallel multi sequence quick select.

---

```
/* Let  $i \in 0 \dots p-1$  and  $p$  denote the local thread index and thread
   count. The number of sequences is equal to  $p$ . The desired
   selection size is  $s$ . */
```

```
 $l_i :=$  start of  $\text{seq}_i$ 
 $u_i :=$  end of  $\text{seq}_i$ 
while  $\sum_{i=0}^{p-1} r_i \neq \text{desired selection sizes}$  do
    pick a thread  $r$  at random
    if  $i = r$  then pick a pivot  $d \in [\text{seq}_i[l_i], \text{seq}_i[u_i]]$ 
     $t_i := \text{binarySearch}(d, [\text{seq}_i[l_i], \text{seq}_i[u_i]])$ 
    if  $\sum_{i=0}^{p-1} t_i < s$  then  $l_i := t_i$ 
    else  $u_i := t_i$ 
end
```

```
/* Thread  $i$  now has the splitter for sequence  $i$  in the variable  $l_i$ .
   */
/* The illustration shows how the range from which the pivot is
   picked decreases with each step. After 3 steps the desired
   selection size is reached in this example. */
```




---

Note that this approach only works if each key is unique. If duplicate keys are allowed it is possible that the correct selection size is never matched. Consider two sequences where one only contains elements with the key 0, the other one elements with the key 1 and our goal is to select  $\frac{1}{4}$  of the elements. The pivot will always be either 0 or 1 and each

step will select  $\frac{1}{2}$  of the elements. To resolve this problem we need two binary searches for each sequence. One search selects the smallest element which is larger than the pivot, the other selects the largest element which is smaller than the pivot. Then we can stop iterating as soon as the target size lies between the sum of the lower ranks and the sum of the upper ranks. A prefix sum over the differences between the upper and lower ranks (i.e. the number of keys equal to the pivot in each sequence) is used in a last step to adjust the selection size: starting with the first sequence, the keys in a sequence equal to the pivot are added to the selection until the desired size is reached. This ensures consistent selections and enables stable merging. Algorithm 3 depicts the modified procedure. We investigate several common strategies for picking the pivot element (e.g. median of three). In our experiments they do not result in significant improvements over picking the pivot at random.

---

**Algorithm 3:** Parallel multi sequence quick select for inputs with duplicate keys.

---

```

/* Let  $i \in 0 \dots p-1$  and  $p$  denote the local thread index and thread
   count. The number of sequences is equal to  $p$ . The desired
   selection size is  $s$ . */

 $l_i := \text{start of seq}_i$ 
 $u_i := \text{end of seq}_i$ 
while  $\sum_{i=0}^{p-1} t_i \leq s$  or  $\sum_{i=0}^{p-1} b_i \geq s$  do
    pick a thread  $r$  at random
    if  $i = r$  then pick a pivot  $d \in [\text{seq}_i[l_i], \text{seq}_i[u_i]]$ 
    /* Rank the pivot from above and below. */

     $t_i := \text{binarySearchGreaterThan}(d, [\text{seq}_i[l_i], \text{seq}_i[u_i]])$ 
     $b_i := \text{binarySearchSmallerThan}(d, [\text{seq}_i[l_i], \text{seq}_i[t_i]])$ 
    if  $\sum_{i=0}^{p-1} t_i > s$  then  $u_i := b_i$ 
    else if  $\sum_{i=0}^{p-1} b_i < s$  then  $l_i := b_i$ 
end

/* Now deal with the keys equal to the last pivot. */
 $\text{gap} := s - \sum_{i=0}^{p-1} b_i$ 
/* The number of keys equal to the pivot in sequence  $i$ . */
 $x_i := t_i - b_i$ 
 $\text{scan}(x_0 \dots x_{p-1})$ 
if  $x_{i+1} \leq \text{gap}$  then  $l_i := u_i$ 
else if  $\text{gap} - x_i > 0$  then  $l_i := l_i + \text{gap} - x_i$ 
/* Thread  $i$  now has the splitter for sequence  $i$  in the variable  $l_i$ .
   */

```

---



### 5.3 Parallel merging

As described in Section 5.1, routines which allow the threads of a thread block to work cooperatively on merging two or more sequences are needed. Parallel merging algorithms which are work-optimal, step-optimal and practical still pose an open problem. Implementations of Coles optimal parallel merge sort algorithm [11] have so far been unable to compete with simpler but asymptotically less efficient algorithms. We implement the well-established odd-even and bitonic merging networks [5], which have a work complexity of  $\mathcal{O}(n \log(n))$ . In addition we implement hybrid algorithms which combine merging networks of a fixed size with two-way and four-way merging. These are based on the ideas of Inoue et al. [21] and can merge sequences with  $\mathcal{O}(n)$  work. An additional work factor, logarithmic in the number of threads, is the price for the exposed parallelism.

#### 5.3.1 Comparator networks




Figure 5: A comparator network with two stages. Note that this is not a sorting network.

We provide a brief introduction to comparator networks. A detailed explanation is given among others by Lang [24]. Comparator networks are sequences of comparator stages. In turn, comparator stages are sets of disjunct comparator elements. A comparator  $C_{i,j}$ , where  $i < j$  and  $i, j \in [1, n]$  for a network of size  $n$ , compares the input elements  $e_i, e_j$  and exchanges them if  $e_i > e_j$ . Comparator networks which have the property to sort their input are termed *sorting networks*.




Figure 6: A bitonic sorting network of size 8. The input is sorted in 6 steps (from left to right).

Sorting networks – and analogously also merging networks – have the property that all comparators in a stage can be executed in parallel. Further, they are data-oblivious and operate in-place. These properties make them ideal for implementation in SIMD

programming models, and also for hardware implementation. The two oldest but still most practical rules for constructing sorting and merging networks are the odd-even merge network and the bitonic merge network [5].

### 5.3.2 Odd-even and bitonic merge

Despite their venerable age of more than 40 years, the bitonic and odd-even merging network are arguably still the most practical parallel data-oblivious merging algorithms. They have a step complexity of  $\mathcal{O}(\log(n))$  and a work complexity of  $\mathcal{O}(n \log(n))$ . Closely related to these merging networks are sorting networks, which can be constructed among others by concatenating  $\log(n)$  merging networks to form a network that sort inputs of size  $n$  in  $\mathcal{O}(\log(n)^2)$  steps. For some specific input sizes more efficient sorting networks are known. It remains an open question whether algorithms for constructing arbitrarily sized sorting networks which have a better asymptotic work complexity and practical constant factors do exist. The AKS network [3] achieves optimal step and work complexity – at the cost of *huge* constant factors. A comprehensive survey of sorting network research is given by Knuth [23].

When choosing between different merging networks we have to differentiate between when we need it to operate on the local memory of a thread block, or where several thread blocks merge cooperatively via the global memory of the GPU. In the first case, both bitonic merging and odd-even merging are viable. Odd-even merging can also be implemented so that it provides a *stable* merge – the stable variant has larger constant factors.




Figure 7: A blocked and cyclic mapping between 8 memory addresses and 2 thread blocks, each consisting of 4 threads.

If global memory access is involved and stable merging is not needed then bitonic merging becomes more attractive. By alternating between two different data layouts it is possible to reduce the amount of global memory access by a factor logarithmic in the size of the local multiprocessor memory, see the work of Ionescu and Schauser [22] for a proof. Let  $i$  denote the index of thread  $t_i$  in the thread block  $b_j$  which lies at position  $j$  in the sequence of  $J$  thread blocks, each of size  $I$ . In the *blocked* layout, each thread processes the elements  $i \cdot j$  and  $(i \cdot j) + I$ . The next  $\log(I)$  steps are executed in the local multiprocessor memory. In the *cyclic* layout the threads of block 1 process the elements  $0, J, \dots, I \cdot (J - 1)$ , the threads of block 2 process the elements  $J + 1, \dots, I \cdot (J - 1) + 1$  and so on. See Figure 7 for an illustration. I.e. the cyclic layout involves strided memory access. Since the local memories are too small to allow for more than 16 steps per layout

switch – the number of steps necessary to outweigh the higher cost of strided access – we always use the blocked layout and perform half of the steps in global memory.

Efficient implementations of merging and sorting networks assume input sizes that are powers of 2. To deal with arbitrary input sizes the input is either padded with sentinel elements to the next larger power of 2, or the network is modified so that its last elements perform a NOOP. Padding incurs some additional index calculations, if multiple bitonic merges are performed in parallel. For an input of size  $n$  it also requires up to  $\frac{n}{2}$  additional memory. Therefore, the approach that does not need padding is preferable.

A further advantage of bitonic merging lies in the fact that conceptually it does not merge two sorted sequences into one sorted sequence. It rather transforms a so-called *bitonic* sequence into a sorted sequence. A bitonic sequence is a concatenation of an ascending sequence and a descending sequence, or vice versa. Since only the total length of the bitonic sequence matters to the network, the overhead for handling the case where both input sequences are not powers of 2 in length is lower for bitonic merging than it is for odd-even merging.

Bitonic merging with local memory usage is one of the strategies we take into closer consideration for our merge primitive. We analyze our implementation in the way that is outlined in Section 3.3. To merge  $n$  elements our implementation uses  $n \frac{\log(n)}{2} + n$  global memory read operations and the same number of write operations. The number of compute instructions we count is  $14n \log(n) + 12n$  – plus a small constant overhead. On the Tesla c1060 and the GTX480 device it is therefore generally bounded by computation.

### 5.3.3 Parallel ranking

Satish et al. [34] use parallel ranking for merging two sequences in parallel: each thread determines the rank of its element in the other sequence via binary search. Asymptotically, parallel ranking has the same complexity as the previously described merging networks. It requires, however, only 2 synchronization barriers instead of  $\log(n)$  and only 1 write operation per element instead of  $\log(n)$ . Further, parallel ranking also is a stable merging algorithm. Its main disadvantage is an irregular memory access pattern and greater memory consumption. Merging networks *exchange* elements and this requires just one register as temporary storage for the exchange operation. For parallel ranking we need two registers for each element a thread processes: one to buffer the element, the other to memorize its rank. I.e. merging networks require constant register space per thread, parallel ranking requires register space linear in the number of elements per thread. Parallel ranking performs notably better on Nvidia Fermi based (e.g. Geforce GTX480) GPUs than on older Nvidia GPUs; Fermi has more flexible broadcasting for local memory, which makes parallel binary searches more efficient.

### 5.3.4 Block-wise merging

To our knowledge Inoue et al. [21] were the first to design a merging algorithm which exposes parallelism by incorporating a merging network into a sequential merging routine. Algorithm 4 follows their idea. We merge *blocks* that have the width  $b$ , which is also the size of a thread block. The resulting work complexity is  $\log(b)n$ .

For both input sequences one block of size  $b$  is buffered; each thread keeps one element of both sequences in registers. Merging is done in a second buffer of size  $2b$  in the

multiprocessors local memory. One step of the merging routine works as follows: thread 0 compares the elements it has stored in registers and writes the result of the comparison to a variable in the shared local memory. Depending on the result of the comparison, each thread then pushes its element from one of the sequences into the second buffer. Effectively, the block from the sequence where the first element is *smaller* is moved to the merge buffer. If the first elements are equal to each other, the block from the sequence which comes first in the input is chosen. The buffer that resides in registers is re-filled from that sequence. Now the merge buffer contains two sorted sequences of length  $b$  – in each step we re-fill the merge buffers first  $b$  positions, the positions  $b \dots 2b$  are filled once at the start in the same manner. E.g. by using a merging network the two sequences are then merged; the first  $b$  positions of the merged result are written to the output. Obviously, the last  $b$  positions of the merge buffer still hold a sorted sequence, which is merged with a new block in the next step. We repeat this process until both input sequences are exhausted. A proof of correctness is given in Appendix B.

Stable merging requires a modification: in the beginning the first  $b$  positions are filled, and after each step the elements at the positions  $b \dots 2b$  are moved to the first  $b$  positions, and then the positions  $b \dots 2b$  are re-filled. This ensures that equal keys do keep their order. Of course the actual merging routine must be stable as well.

---

**Algorithm 4:** Block-wise merging.

---


```

/* Let  $i \in 0 \dots p-1$  and  $p$  denote the local thread index and thread
   count. The merge block size is equal to  $p$ . */

for  $o := i$  ;  $o < |input_A| + |input_B|$  ;  $o := o + p$  do
/* Use a merging network of size  $2 \times$  block size  $p$ . */
merge( $buffer_0 \dots buffer_{2 \cdot p}$ )
output[o] = buffer[i]
/* Thread 0 determines which block to load next. */
if  $i = 0$  then  $aIsNext := nextA \leq nextB$ 
/* Put the next element from  $A$  or  $B$  into the buffer. Advance
   position in  $A$  or  $B$  by  $p$  and load the next element. */
if  $aIsNext$  then  $buffer[i] := getNext(input_A)$ 
else  $buffer[i] := getNext(input_B)$ 

```

---



It turns out that the optimum block size should be the SIMD width of the hardware. Because then the increased work complexity is of no consequence. This disregards two aspects of the algorithm: first, per block one additional comparison that is made by a single thread is required. It follows that the amount of sequential work decreases for larger block sizes. More important, the block size determines how many parallel merging processes we need to saturate the multiprocessors of the GPU. What we gain here in efficiency would be outweighed by a more expensive multi sequence selection step.

Referring to Section 3.5 we say that block-wise merging is SIMD-conscious, and that multi sequence selection is multiprocessor-conscious. A merging routine based on multi sequence selection and block-wise merging is therefore both SIMD- and multiprocessor-conscious.

An important detail of our implementation is that we fill the buffer registers with *sentinels* when a sequence is exhausted. A sentinel is a dummy element which is greater than any element in the input data. The implementation is greatly simplified by using sentinels. The merging network also need not be able to deal with inputs that are not powers of 2. This decreases the key range by 1, however, for stable merging; if there are keys with the same value as the sentinel element the algorithm may write out sentinels. For non-stable merging we end up using the bitonic merge and for stable merging we employ parallel ranking, see Section 5.5.4.

Again, we analyze the work complexity of our actual implementation. Per element our two-way merging implementation takes  $12 + 2 \cdot \text{SIMD-width/block-size}$  instructions outside the merging network, plus a small constant amount of work. Since one global read and write is required we are left with  $6 + \text{SIMD-width/block-size}$  compute instructions per memory operation, without even considering the work of the merging network. This implies that our two-way merging is compute-bounded.

### 5.3.5 Pipelined merging

We explore pipelined merging as a means of decreasing the amount of global memory access and a possible way to increase work efficiency. Pipelined merging is commonly used in models with hierarchical memory [16]. Conceptually, one builds a binary tree of two-way mergers – see Figure 8. Between the levels of the tree buffers are placed which reside in fast local memory (i.e. local multiprocessor memory for GPUs) and all mergers except for the ones in the first level are fed from buffers. A tree of  $2k - 1$  two-way mergers can merge  $k$  sequences of combined length  $n$  with  $\mathcal{O}(n)$  accesses to global memory. Successive two-way merging would incur  $\mathcal{O}(\log(k)n)$  expensive memory accesses. The limiting factor for implementing a pipelined technique on GPUs is the available amount of local memory.




Figure 8: By using a tree of two-way mergers the merging process is pipelined. Buffers between the stages of the tree then serve to replace expensive access to global memory by access to local memory.

To evaluate pipelined merging on GPUs we implement the most basic variant: four-way merging. Three of our two-way block-wise mergers are combined into a tree, with buffers between the two levels of the tree. In each step all threads of the thread block operate one of the mergers on the first level of the tree and then on the head of the tree. Following the method of analysis we explain in Section 3.3, four-way merge incurs  $21 + 4 \cdot \text{SIMD-width}/\text{block-size}$  instructions outside the merging networks per element. In our simplified model it is therefore *slightly* more efficient than two-way merging. In comparison to two-way merging, however, our four-way implementation has to place more variables in local memory instead of registers, requires more registers per thread and also more buffer space in local memory.

### 5.3.6 Other parallel merging approaches

Another way of  $k$ -way merging uses a priority queue for selecting the next smallest element from  $k$  sequences. The priority queue is commonly implemented as a tournament tree – an inherently sequential data structure. As an alternative one could use a parallel reduction instead of a tournament tree. However, the parallel reduction adds an additional work factor of  $k$ . We implement both and find them to deliver abysmal performance in practice, which is unsurprising. The tournament tree of size  $k$  needs to be stored in local memory to actually reduce the number of accesses to global memory in comparison to two-way merging. Also, to allow for efficient reading from global memory, an input buffer of 16 elements per sequence must be used. This means only very few tournament tree mergers can run on a multiprocessor – and only one thread can be used per tournament tree. The parallel reduction fares just as bad due to its suboptimal work complexity.

Coles merge sort [11] achieves optimal step complexity and work complexity in the PRAM model. Essentially it uses selection recursively, combined with pipelined merging. Ding et al. [13] implemented a merging primitive based on Coles idea. The method proposed by Hagerup and Rüb [18] also merges recursively; Satish et al. [34] derived a GPU merging algorithm from it.

## 5.4 Sorting by merging

We use our merging routine to construct a sorting algorithm. Sorting by merging has several advantageous properties. In contrast to distribution-based algorithms, merge sort incurs strictly sequential memory write access. Merge sort further allows to distribute work evenly among multiple processors. Lastly, its performance is largely independent of the input key distribution.

Having a merging primitive available, the only additional complication is producing the initial set of sorted sequences. We compare four different sorting routines: bitonic merge sort, quicksort, two-way block-wise merge sort and four-way block-wise merge sort. All of these can be run with  $k$  thread blocks to create  $k$  sorted sequences concurrently. The merge based sorting routines follow the strategies we describe in this section. See Appendix C for a description of the parallel quicksort algorithm we use.

## 5.5 Experimental results

To ensure that our results are reproducible we describe our hardware setup and our methods for generating input data.

### 5.5.1 Hardware setup

Our test hardware consists of an Intel Q6600 CPU with 8 GiB RAM. We use a Gigabyte Geforce GTX480. The system runs OpenSUSE 11.1 64 bit Linux. Everything is compiled with GCC 4.3.4. The Nvidia CUDA 3.0 compiler and runtime environment are used for all benchmarks. All performance numbers exclude the time for transferring input to the GPU memory and back to the system main memory. Instead of running time our charts show the *number of elements processed per time*, respectively for selecting, merging and sorting.

### 5.5.2 Input data

For our analysis we use a commonly accepted set of distributions motivated and described in [19]. We describe the distributions with the parameters that are used to create 32 bit keys. 64 bit keys are generated analogously.

**Uniform.** A uniformly distributed random input in the range  $[0, 2^{32} - 1]$ .

**Gaussian.** A gaussian distributed random input approximated by setting each value to an average of 4 random values.

**Bucket Sorted.** For  $p \in \mathbb{N}$ , the input of size  $n$  is split into  $p$  blocks, such that the first  $\frac{n}{p^2}$  elements in each of them are random numbers in  $[0, \frac{2^{32}}{p} - 1]$ , the second  $\frac{n}{p^2}$  elements in  $[\frac{2^{32}}{p}, \frac{2 \cdot 2^{32}}{p} - 1]$ , and the last  $\frac{n}{p^2}$  elements in  $[\frac{(p-1) \cdot 2^{32}}{p}, 2^{32}]$ .

**Staggered.** For  $p \in \mathbb{N}$ , the input of size  $n$  is split into  $p$  blocks such that if the block index is  $i \leq \frac{p}{2}$  all its  $\frac{n}{p}$  elements are set to a random number in  $[(2i-1)\frac{2^{31}}{p}, (2i)\frac{2^{31}}{p-1}]$ .

**g-Group.** For  $p \in \mathbb{N}$ , the input of size  $n$  is split into  $p$  blocks. Consecutive blocks form groups consisting of  $g \in \mathbb{N}$  blocks. In group  $j$  the first  $\frac{n}{pg}$  elements are random numbers between  $\left(\left(\left((j-1)g + \frac{p}{2} - 1\right) \pmod{p}\right) + 1\right) \frac{2^{31}}{p}$  and  $\left(\left(\left((j-1)g + \frac{p}{2}\right) \pmod{p}\right) + 1\right) \frac{2^{31}}{p} - 1$ . The second  $\frac{n}{pg}$  elements are set to be random numbers between  $\left(\left(\left((j-1)g + \frac{p}{2}\right) \pmod{p}\right) + 1\right) \frac{2^{31}}{p}$  and  $\left(\left(\left((j-1)g + \frac{p}{2} + 1\right) \pmod{p}\right) + 1\right) \frac{2^{31}}{p} - 1$ , and so forth.

**Deterministic Duplicates.** For  $p \in \mathbb{N}$ , the input of size  $n$  is split into  $p$  blocks, such that the elements of the first  $\frac{p}{2}$  blocks are set to  $\log n$ , the elements of the second  $\frac{p}{4}$  processors are set to  $\log(\frac{n}{2})$ , and so forth.

**Randomized Duplicates.** For  $p \in \mathbb{N}$ , the input of size  $n$  is split into  $p$  blocks. For each block an array  $T$  is filled with a constant number  $r$  of random values set between 0 and  $r - 1$  - the sum of these values is  $S$ . The first  $\frac{T[1]}{S} \times \frac{n}{p}$  elements of the input are set to a random value between 0 and  $r - 1$ , the next  $\frac{T[2]}{S} \times \frac{n}{p}$  elements of the input are set to another random value between 0 and  $r - 1$ , and so forth.

**Sorted.** A uniformly distributed random input in the range  $[0, 2^{32} - 1]$ , which has been sorted.

Our source for uniform random values is the Mersenne Twister [26]. The parameters for the distributions are set to:  $p = 240$ ,  $g = 8$  and  $r = 32$ . Thus,  $p$  corresponds to the number of scalar processors on an G200-based Tesla c1060 and  $g$  to the SIMD width of its multiprocessors. For  $r$  we take the value that was chosen in the original work which introduced these input distributions. While we perform our final experiments on a Fermi-based Geforce GTX480, which has a different number of scalar processors and a different SIMD width than G200 has, we stay with the parameters we used in our previous work to ensure the results are comparable to each other.

In this section we explore the parameter space for the previously explained algorithms. Then we assemble them into primitives for merging of key-only inputs, stable and non-stable merging of key-value inputs and for sorting.

Unless stated otherwise we use uniformly distributed 32 bit keys as input. Input for merging and selection is generated from  $k \cdot m$  keys taken from the chosen input distribution. These  $k \cdot m$  keys are split evenly into  $k$  sequences which are then sorted. Note that we only give the results for a meaningful selection of benchmarks. Covering all or most of the possible combinations of parameters would neither be possible due to lack of space nor would it be insightful.

### 5.5.3 Multi sequence selection

At first we take a look at how both multi sequence selection (also called *multiway selection*) routines perform for different numbers and lengths of sequences. The number of splitters is set to 60 – a multiple of the multiprocessor count. The GPU implementation of Vitaly Osipovs algorithm has a work complexity of  $\mathcal{O}(k^2 \log(m))$  and a step complexity of  $\mathcal{O}(k \log(k) \log(fm))$ . However, it merely needs  $\mathcal{O}(k \log(m))$  accesses to global memory. Our multi sequence quick select implementation has an expected work complexity



of  $\mathcal{O}(k \log(m)^2)$ , an expected step complexity of  $\mathcal{O}(\log(m)^2)$  and it requires expected  $\mathcal{O}(k \log(m)^2)$  accesses to global memory. Since the essential operation of both algorithms is binary search, there is little computation that could hide memory access latency. We assume that Osipovs algorithm outperforms the other for large sequences due to its lower work complexity and bandwidth requirements. The results shown in Figure 9 confirm our assumption only partly: quick select is faster even for sequences of length  $2^{20}$ , if the input consists of 32 or more sequences.

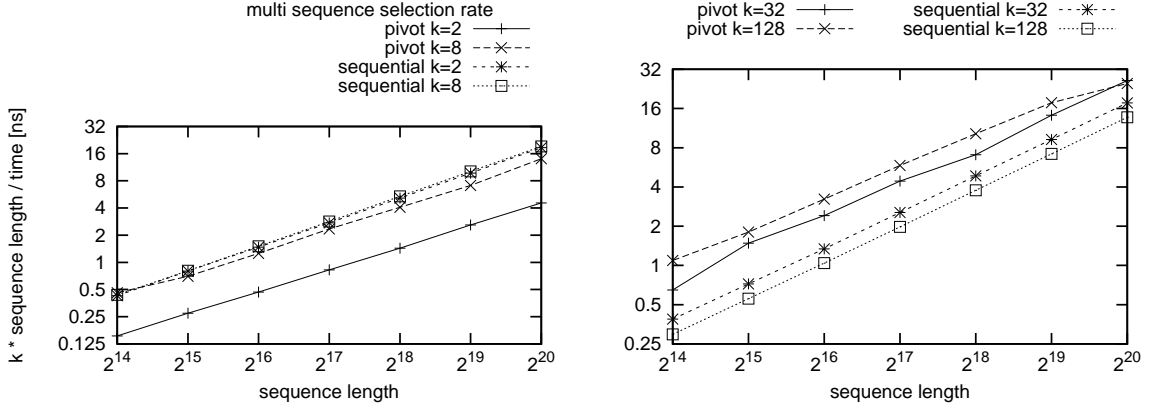


Figure 9: Performance (elements processed per time) of sequential and pivot-based multi sequence selection with varying sequence count  $k$  for different input sizes.

We expect our pivot-based selection to perform differently on input distributions that deviate from the uniform distribution. Osipovs algorithm should be oblivious to the distribution of the input. Here our assumptions are fully confirmed by our experiments, see Figure 10. The performance of multi sequence quick select deteriorates to an impractical level on the staggered distribution. We therefore use Osipovs selection algorithm for our merging primitive.

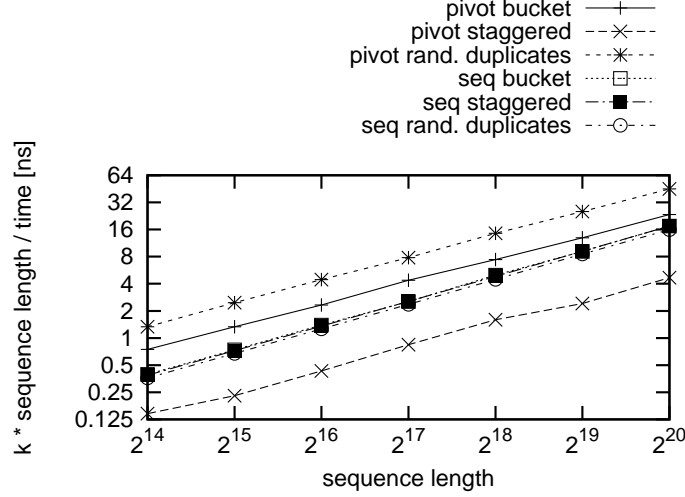


Figure 10: Performance (elements processed per time) of sequential and pivot-based multi sequence selection for different input distributions,  $k$  is set to 32. While the sequential selection algorithm is oblivious to the input distribution, the pivot-based selection shows large variance in performance for different input distributions.

The running time for varying numbers of splitters implies a trade off between the work performed by multi sequence selection and merging. Therefore we will look at it in the context of the whole merging primitive.

#### 5.5.4 Merging networks

Since the bitonic merge network, the odd-even merge network and parallel ranking all have the same asymptotic work and step complexity we expect no large variance in performance between them. We measure their dynamic instruction count for different thread block sizes, in the context of block-wise two-way merging. The instruction counts are generated by subtracting the counts from a skeleton program, which performs the same work but does not execute a merging network, from those of the full program.

At the time of writing the profiling tool for Nvidia GPUs are unable to display instruction counts for Fermi-based cards. We therefore measure them on a G200-based card. We expect that the instruction count on Fermi will be lower for parallel ranking, due to its improved local memory broadcasting capabilities.

Figure 11 shows a comparison between the measured dynamic instruction counts and our estimates. All estimates are based on counting instructions in our actual implementations, as described in Section 3.3. For both merging networks our estimates turn out to be too large. This is no surprise since they assume the worst case, where in each step all elements have to be exchanged. For parallel ranking the situation is reversed: we do not account for memory access conflicts that occur during the concurrent binary searches and thus our instruction counting is too optimistic.

These numbers confirm our hypothesis of block-wise merging being bound by computation. Even the most favorable case (bitonic merge with 32 threads) takes more than 3 times the number of compute instructions that would be necessary to avoid limitation by bandwidth on the Tesla c1060.

algorithm	block size	32	64	128	256
<i>odd-even</i>					
dynamic instructions		59	65.3	74.3	85.3
estimate: $8 + 15(\log(n) - 2)$		53	68	83	98
<i>bitonic</i>					
dynamic instructions		42.9	50.8	59	70
estimate: $13(\log(n) - 1)$		52	65	78	91
<i>ranking</i>					
dynamic instructions		81.8	93.3	106.4	121.1
estimate: $7 + 11(\log(n) - 1)$		51	62	73	84

Figure 11: Per-element instruction count estimates and dynamic measurements of odd-even merge, bitonic merge and parallel ranking for different thread block sizes.

### 5.5.5 Merging

Here we are interested in the trade off between multi sequence selection and merging for each of our three merging routines. Our merging algorithms ought to be more efficient at merging than our multi sequence selection algorithms. Thus we expect the best running time for two-way and four-way when the work is split among just enough merging thread blocks to saturate all multiprocessors of the GPU. For bitonic merging this may look different because of the additional log-factor in its work complexity. We take the average time per merged element for sequences in the range of  $[2^{12}, 2^{20}]$ . The number of merging thread blocks is incremented in steps of 30 – the number of multiprocessors of our GPU. Our implementations are limited to thread block sizes that are powers of 2; therefore we measure performance for 32, 64, 128 and 256 threads per block. As multi sequence selection routine we use Algorithm 1.

The results do match our expectations, see Figure 12. Two-way and four-way merging perform best when there are exactly as many merging thread blocks as can execute on the GPU concurrently. We measure peak performance for two-way merging when there are 120 mergers with 128 threads each. The worse performance for merging thread blocks with 64 and 32 threads is caused by a multiprocessor only executing 8 thread blocks concurrently. Therefore the GTX480s multiprocessors cannot run enough threads when using merging thread blocks with less than 128 threads. On the G200-based Tesla c1060 – which requires fewer threads to saturate its multiprocessors – the best performance is achieved for 64 threads per thread block.

Bitonic merging profits from greater merger counts, since then the sequence parts it has to merge get smaller. Overall, two-way merging shows the best performance.

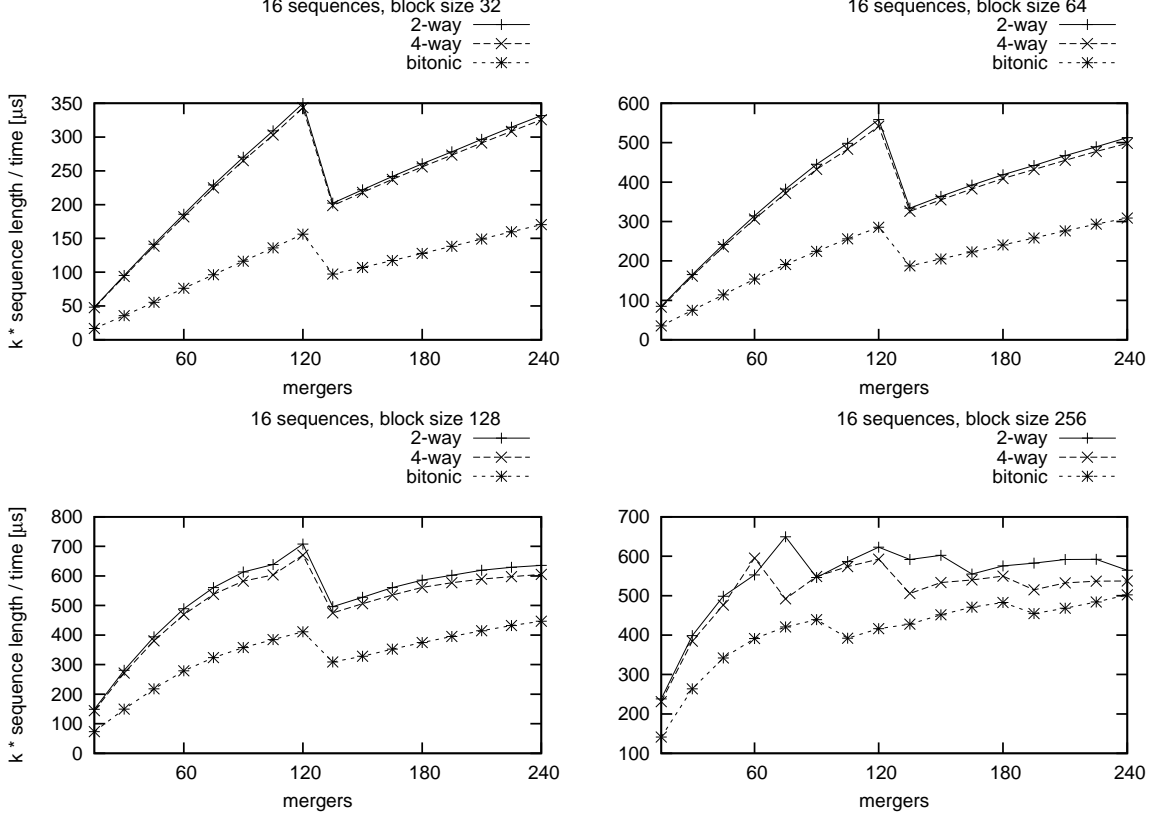


Figure 12: Performance of bitonic, two-way and four-way merging for varying numbers of merging thread blocks. We show the average of the performance for input sequences with lengths in the range of  $[2^{12}, 2^{20}]$ .

To put the performance of our merging primitive into perspective we compare it to the merge step of Thrust merge sort [34]. We assume that a merging primitive based on the merge sort implementation by Ye et al. [40] would perform approximately as well as ours. The implementation by Ding et al. [13] is slower than Thrusts merge, according to their published results. Our primitive uses 120 parallel mergers with 128 threads each. The non-stable version uses a bitonic merging network, the stable one uses parallel ranking. Our advantage over the merge from Thrust is on average 35.2% for the non-stable merge variant and 33.5% for the stable variant. A recent merge sort implementation by Satish et al. [35], based on selection in two sorted sequences, outperforms a sort based on our merging primitive by a factor of 2-3. We have, at the time of writing, no access to the implementation, and therefore do not know how it performs at merging large sequences. But from its sorting performance we expect it to outperform our merging routine to a similar degree.




Figure 13: Comparison between our stable and non-stable key-value merging primitive and the merge step from Thrust merge sort.

### 5.5.6 Sorting

Lastly, we construct a sorting algorithm. We use quicksort, bitonic sort and sorters based on our two-way and four-way merging algorithms to create a set of sorted sequences from the input – one thread block creates one sorted sequence. We find that parallel quicksort outperforms bitonic sort and also sorters based on our two-way and four-way merging algorithm. We merge these sorted sequences with our merging primitive to obtain the sorted result.

See Figure 14 for a comparison between our merge sort, GPU Sample Sort and CUDPPs radix sort. Merge sort is comfortably beaten by GPU Sample Sort. Note that we did not spend time to tune merge sort specifically for small inputs. Our implementations performance for small inputs suffers because we either have to run multi sequence selection on small sequences (where it is inefficient), or we can run only a small number of thread blocks to create the initial sorted sequences, leaving multiprocessors unoccupied.

We assume that fundamental algorithmic improvements would be necessary to make our merge sort strategy competitive. The results of Ye et al. [40] support our argument – despite a considerably higher amount of architecture-specific optimization their merge sort implementation, which follows the same overall strategy, achieves similar results.

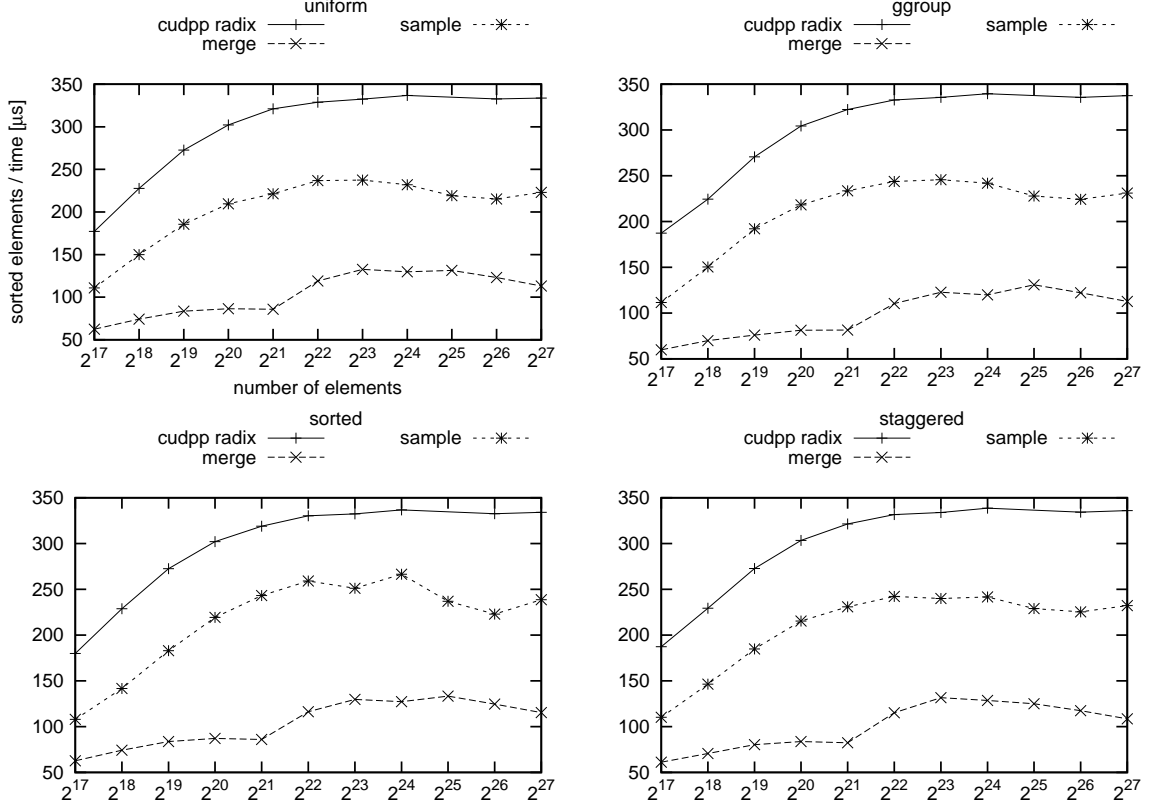


Figure 14: Comparison between quick-multi-selection merge sort, GPU Sample Sort and CUDPPs radix sort for different key distributions.

## 6 Revisiting GPU Sample sort

As we are unable to match GPU Sample Sorts performance with our merge sort variants we attempt to improve GPU Sample Sort, and use it to evaluate the current state of the OpenCL programming language in Section 7.

### 6.1 Algorithm overview

Here we describe our original sample sort algorithm.

Randomized sample sort is considered to be one of the most efficient comparison-based sorting algorithms for distributed memory architectures. Its sequential version is best described in pseudocode, see Algorithm 5. The oversampling factor  $a$  trades off the overhead for sorting the splitters and the accuracy of partitioning.

The splitters partition input elements into  $k$  buckets delimited by successive splitters, where  $k$  is a power of 2. Each bucket can then be sorted recursively, and their concatenation forms the sorted output. If  $M$  is the size of the input when `SmallSort` is applied, the algorithm requires  $\log_k \frac{n}{M}$   $k$ -way distribution passes in expectation until the whole input is split into  $\frac{n}{M}$  buckets. Using a small case sorter with a work complexity of  $\mathcal{O}(n \log n)$  leads to an expected work complexity of  $\mathcal{O}(n \log n)$  overall.

---

**Algorithm 5:** Serial Sample Sort.

---

```

SampleSort( $e = \langle e_1, \dots, e_n \rangle, k$ )
begin
  if  $n < M$  then return SmallSort( $e$ )
  choose a random sample  $S = S_1, S_{2a} \dots, S_{(k-1)a}$  of  $e$ 
  Sort( $S$ )
   $\langle s_0, s_1, \dots, s_k \rangle = \langle -\infty, S_a, S_{2a}, \dots, S_{(k-1)a}, \infty \rangle$ 
  for  $1 \leq i \leq n$  do
    find  $j \in \{1, \dots, k\}$ , such that  $s_{j-1} \leq e_i \leq s_j$ 
    place  $e_i$  in bucket  $b_j$ 
  return Concatenate(SampleSort( $b_1, k$ ), ..., SampleSort( $b_k, k$ ))
end
end

```

---

Since the small buckets are sorted in parallel, load-balancing depends on the variance in size between the buckets, i.e. the quality of the splitter elements. Sufficiently large random samples yield provably good splitters independent of the input distribution [8].

Due to the high cost of global memory accesses on GPUs, multi-way approaches are more promising than two-way: each  $k$ -way distribution step requires  $\mathcal{O}(n)$  memory accesses. Expected  $\log_k \frac{n}{M}$  passes are needed until buckets fit into fast GPU local multi-processor memory of size  $M$ . Thus, we can expect  $\mathcal{O}(n) \log_k \frac{n}{M}$  global memory accesses instead of the  $\mathcal{O}(n) \log_2 \frac{n}{M}$  accesses required by two-way sorting. However, the write access of a  $k$ -way distribution step is not sequential but scattered, reducing the effective memory bandwidth.




Figure 15: An iteration of  $k$ -way distribution.

GPU Sample Sort follows a three level strategy. Initially there is one bucket: the whole input of size  $n$ . If a bucket is larger than a threshold  $M$  (a size below which further  $k$ -way distribution would be inefficient), we apply our  $k$ -way distribution procedure to it. The work of  $k$ -way distribution is mostly data-parallel and can be assigned to as many

as  $\frac{n}{t}$  thread blocks, where each thread block has up to  $t$  threads. When no buckets larger than  $M$  are left, the buckets are sorted concurrently with quicksort. Quicksort switches to odd-even merge sort [5] for buckets that fit into a multiprocessors fast local memory. In the quicksort step we map one thread block to one bucket. This looks like a problem, as the amount of work may vary significantly for different thread blocks. In practice, however, there are typically a lot more thread blocks than multiprocessors and we achieve sufficient load-balancing.

GPU Sample-Sort's distribution pass, when the bucket size exceeds a fixed size  $M$ , can be described in 4 steps corresponding to individual kernel launches, see Figure 15.

**Step 1.** Choose splitters as in Algorithm 5.

**Step 2.** Each of the  $p$  thread blocks computes the bucket indices for all elements in its tile, counts the number of elements in each bucket and stores this per-block  $k$ -entry histogram in global memory.

**Step 3.** Perform a prefix sum over the  $k \times p$  histogram tables stored in a column-major order to compute global bucket offsets in the output.

**Step 4.** For its elements each of the  $p$  thread blocks computes their local offsets in the buckets. Finally it stores elements at their proper output positions, using the global offsets computed in the previous step.

## 6.2 Algorithm implementation

To reiterate our strategy: buckets larger than some threshold are split via  $k$ -way distribution. Quicksort splits them up further until they fit into the local memory of a thread block. Then we apply odd-even merge sort. The first two levels of sorting require temporary space in GPU memory of the size of the input. Additional global memory holds the prefix sum histogram tables. For keeping track of buckets during the distribution pass a small amount of CPU memory is required. The size of the prefix sum tables is bound by a constant and all other temporary storage we use is small enough to be neglected. Thus, we say our implementation uses space twice the size of the input.

**Step 1.** We take a random sample  $S$  of  $a \cdot k$  input elements. A linear congruential generator is used to synthesize pseudo-random numbers on the GPU. The generator gets its seed from a Mersenne Twister which runs on the host CPU. We sort  $S$  and place each  $a$ -th element of it in the array of splitters  $bt$  such that they form a complete binary search tree with  $bt[1] = s_{k/2}$  as the root. The left child of  $b[j]$  is placed at the position  $2j$  and the right one at  $2j + 1$ , see Algorithm 6. If we find that all keys in  $S$  are equal, we memorize this observation for the other distribution steps. A problem is how to sort  $S$  efficiently on the GPU, as in practice it only contains a few thousand elements. We simply use a single thread block for sorting  $S$ . We run the step in parallel for multiple buckets.

**Step 2.** Now we split the work among multiple thread blocks: blocks are assigned to equal-sized parts (tiles) of the bucket. The previously created binary search tree is first copied into the GPUs cached read-only memory, then loaded into local memory by each




thread block. We adopt a tree traversal method that was previously used to prevent branch mispredictions on commodity CPUs possessing predicated instructions [33], refer to Algorithm 6. GPUs support predicated instructions as well, see Section 3.2 for details.

---

**Algorithm 6:** Serial search tree traversal.

---



$$bt = \langle s_{k/2}, s_{k/4}, s_{3k/4}, s_{k/8}, s_{3k/8}, s_{5k/8}, s_{7k/8} \dots \rangle$$

**TraverseTree**( $e_i$ )

**begin**

$j := 1$

**repeat**  $\log k$  times  $j := 2j + (e_i > bt[j])$  // Left or right child?

$j := j - k + 1$  // Bucket index

**end**

---

Each thread loads an element into a register and traverses the binary search tree, storing its bucket index in local memory. Then, again in local memory, a counter corresponding to that bucket index is incremented. The counters can be replicated, increasing parallelism at the expense of local memory space. This procedure is repeated until the thread blocks whole tile has been processed. The local bucket sizes are written to global memory, forming a  $k \times p$  table –  $p$  being the number of thread blocks.

**Step 3.** Performing a prefix sum over bucket size table yields global bucket offsets for the thread blocks. As efficient prefix sum implementations for GPUs are available we do not build our own. Instead, we take the one from the Thrust library [34].

**Step 4.** We recompute the elements bucket indices and scatter them, incrementing offset counters for each bucket in local memory. An elements output position is given by its buckets global offset plus that buckets current local counter value. Atomic operations are used for incrementing the counters.

We still have to consider the previously mentioned degenerate case: all keys in a sample may be equal. If at the same time their value is greater or equal than that of any other key in the input, the whole input will be distributed to a single bucket. As a result we could recurse infinitely on a bucket, failing to ever split it into smaller buckets. Our solution is to do a partitioning step as in quicksort when we encounter a degenerated sample. We distinguish between elements that are lesser, equal or greater

than the samples key. Buckets containing elements equal to a samples key need no further sorting, obviously. Additionally – if a degenerated sample has been detected – two parallel reductions are performed on the bucket, determining if the bucket contains any keys with unequal values. Step 4 can be skipped if all keys are equal. This distribution-dependent behavior works to our advantage. For input distributions with low entropy GPU Sample Sort can skip superfluous comparisons.

Note that a concise way of avoiding the problem – used for sample sort on other types of machines [8] – is to tag each key with its position in the input. Thereby one guarantees that no duplicate keys exist. The high cost of memory operations and the small amount of local multiprocessor memory make this approach infeasible on GPUs.

This concludes the description of our  $k$ -way distribution procedure. The number of thread blocks can be chosen freely: it allows for a trade-off between the amount of exposed parallelism on the one hand, and an increased amount of work per thread on the other hand. The space required for the global offset table and the cost of the prefix sum also depend on it. Therefore, the  $k$ -way distribution procedure is weakly multiprocessor-conscious (see Section 3.5): the number of thread blocks affects the required amount of work, but this work only makes up a fraction of the total work. It is, however, SIMD-oblivious.

**Sorting buckets.** We delay the sorting of buckets until the whole input is partitioned into buckets of size at most  $M$ . Since the number of buckets grows with the input size, it is larger than the number of processors most of the time. Therefore, we can use a single thread block per bucket without sacrificing exploitable parallelism. To improve load-balancing we schedule buckets for sorting ordered by size. For buckets that are too large to fit into local memory, we use parallel quicksort to split them up further.

For sequences that fit into local memory, we use the odd-even merge sorting network [5]. In our experiments we found it to be faster than the bitonic sorting network and other approaches like a parallel merge sort.

## 6.3 Exploring bottlenecks

We discuss several bottlenecks of GPU Sample Sort and identify aspects which can be improved without changing the overall structure of the algorithm. Further, we give performance numbers for our implementations on the Geforce GTX480. Refer to Section 5.5 for a complete description of our experimental setup.

### 6.3.1 Sorting networks

The base case sorting with odd-even merge sort makes up approximately 40% of the overall running time. However, searching for alternatives to the classic (bitonic and odd-even merge) sorting networks is a deep topic in itself and outside the scope of this work. Possible candidates include quicksort, merge sort, and also other sorting networks which have not yet been considered for implementation in the PRAM model.

### 6.3.2 Second level sorting

With  $M$  as the bucket size below which the second level sorter is used, the second level sorters impact on the running time is greatest when  $M \cdot k^m$ ,  $m \in \mathbb{N}$ , is approximately equal to the input size; then it contributes the most to the sorting process. This shows in our original implementation: its performance decreases for the input sizes where the second level quicksort has to sort larger buckets.

We explore two-way merging, four-way merging, and randomized sample sort as alternatives to quicksort. Merge sort seems attractive because of its sequential memory access pattern, and because the base case sorting can be performed on tiles that are powers of two in size – sorting network implementations are more efficient for inputs that have a size that is a power of 2. In practice we do find that the dynamic instruction count of our merge sort implementations is at least 40% larger than that of quicksort, and that they exhibit worse performance. The advantage of sample sort and quicksort is their synergy with the first level of the sorting process: each  $k$ -way distribution step decreases the entropy within the resulting buckets, and distribution-based sorting gains efficiency for inputs with low entropy. Further, both sample sort and quicksort are SIMD-oblivious. We implement a combination of quicksort and sample sort, which is faster and scales better with the input size than our original implementation, which solely used quicksort.  $M$  is set to  $2^{18}$ , the second level sort uses  $k$ -way distribution with  $k = 32$ , and quicksort for buckets smaller than 12,000. Since the input size is bounded a small oversampling factor suffices (we use  $a = 6$ ); consequently we only have small samples and can sort them directly with odd-even merge sort. Apart from that the second level sample sort implementation is based on the algorithm described in Section 6.1. We show the performance of the different sorting implementations in Figure 16.

### 6.3.3 $k$ -way distribution

The primary work in a  $k$ -way distribution step is done in the kernel that determines the local bucket sizes and in the kernel that scatters the data to the buckets; the prefix sum which is performed between these two kernels has a negligible impact on the running time. Traversal of the binary search tree, atomic increments of the bucket size counters and scattered writes are the costly parts of  $k$ -way distribution.

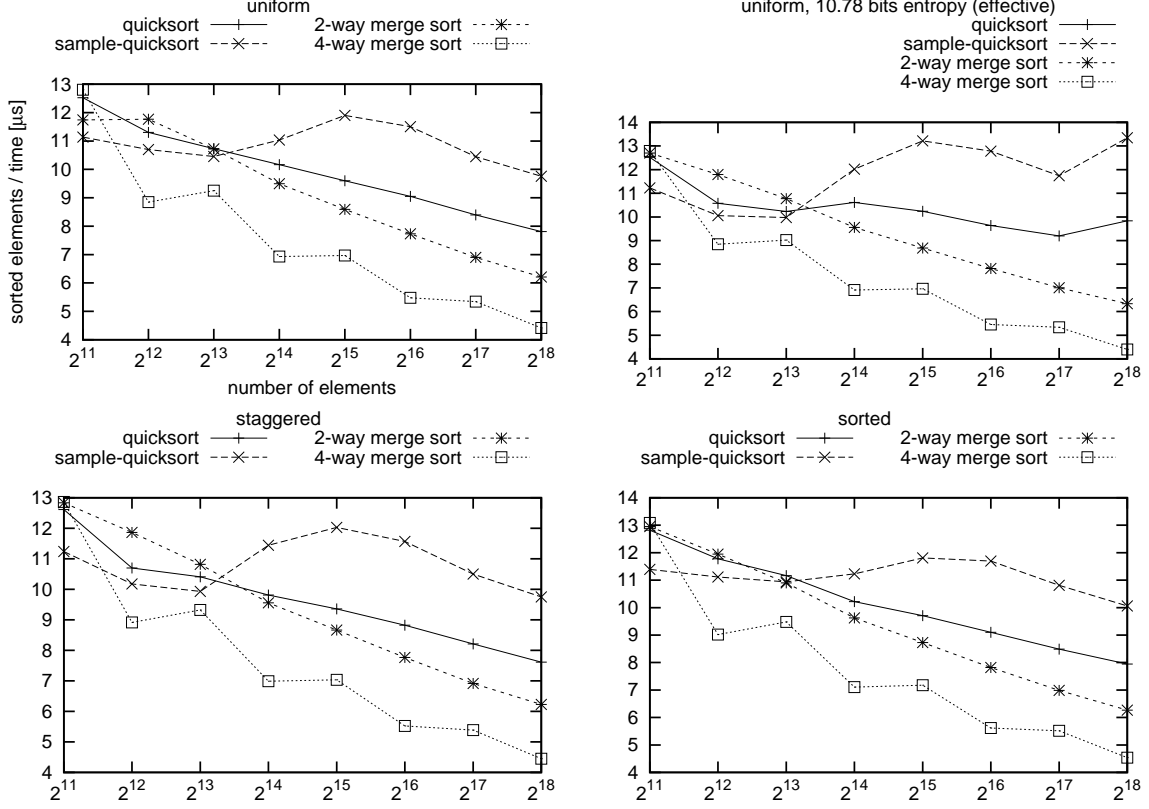


Figure 16: A comparison between the performance of thread block quicksort, sample sort, two-way merge sort and four-way merge sort on the Geforce GTX480. Only a single thread block is run on the device for this benchmark. Quicksort and sample sort are more efficient for low entropy inputs. We lower the entropy by combining multiple keys with a bit-wise and-operation. Combining 4 32 bit keys in this way results in an effective entropy of 10.78 bits. Each sorter uses odd-even merge sort as base case sorter for input sizes smaller than 2049.

In our original implementation traversing the binary search tree is already very efficient; the number of performed comparisons is optimal, no branching occurs and apart from that a traversal step performs only one read access to local memory, an addition and a bit-shift operation on data that lies in a register. The only aspect that seems to be not optimal is the read access to local memory (the binary search tree is stored in local memory) because there may be access conflicts, which lead to serialized access. Nvidias Fermi architecture, however, has local memory with broadcasting capabilities. There are only memory bank conflicts if threads concurrently access different words that belong to the same memory bank. With 32 memory banks and a word width of 32 bit, the first 6 steps of the search tree traversal are conflict-free. The seventh step exhibits two-way conflicts, the eighth exhibits four-way conflicts and so on. We use a binary search tree with 7 levels and therefore only pay one additional instruction in the last step of the traversal because of conflicts. With a larger tree or with keys longer than 32 bit the number of conflicts grows, but the overhead is arguably still smaller than it would be if we would e.g. replicate the binary search tree and use a more complicated indexing

scheme.




Figure 17: The memory layout of the binary search tree for a GPU with 4 memory banks and  $k = 16$ . Only the last level of the tree contains more than one word in each memory bank and will therefore exhibit two-way bank conflicts on an architecture with broadcasting capabilities like Nvidias Fermi.

Since we share bucket size counters between multiple threads, atomic operations are required to increment them (or explicit serialization). The reason is that there is too little local memory and register space available to have separate counters for each thread. For input keys that are distributed in a way that consecutive keys do go into different buckets, there are no conflicts when incrementing the counters. The worst case occurs if all keys processed by the same SIMD group go to the same bucket. Replicating the bucket counters as many times as possible within the limits of the local memory size, so that fewer threads share counters, improves the worst case behavior.

But using more counters has a price. A higher memory consumption per thread lowers the number of threads that can run on a multiprocessor, which gives the multiprocessor a smaller chance to hide memory access latencies. Further, it also decreases the locality of write access in the scattering kernel as there are more possible write locations. The locality of write access (and thus the effective bandwidth) is also highest when there are the most atomic increment conflicts. As a result, depending on the distribution of the input the kernels shift from being computationally bound to being bandwidth-bound.

We use a different number of threads per bucket counters for step 3 (bucket counting) and step 4 (scattering) of the  $k$ -way distribution routine, refer to Section 6.2. Obviously the scattering kernel cannot use fewer threads per counter as the bucket counting kernel, but it can use more threads per counter during scattering. The bucket finding kernel only executes one sequential read per search tree traversal and atomic increment, making it more computationally bound than the scattering kernel. In contrast, the scattering kernel executes an additional scattered write, whose average cost increases with a larger number of bucket counters. Our implementation uses 8 counters per bucket per thread block in the bucket finding kernel, and 4 in the scattering kernel. This provides a balance between a better worst case performance and a lower best case performance.

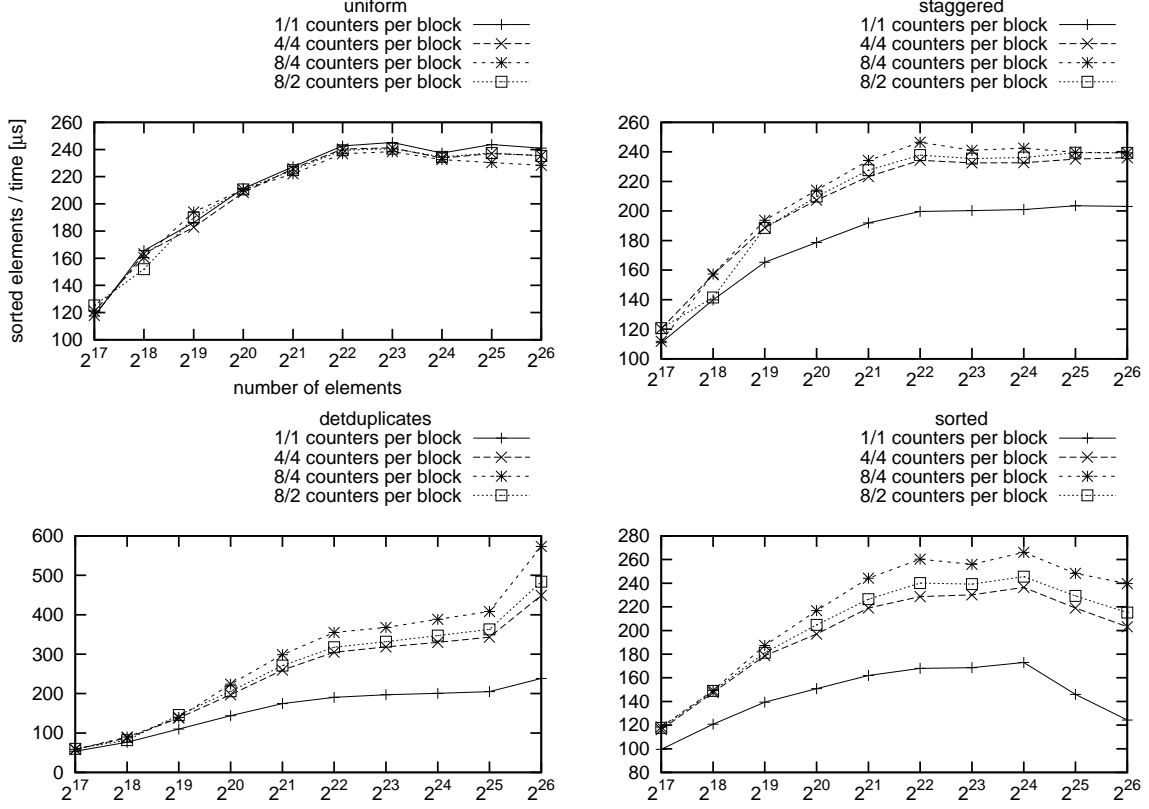


Figure 18: Performance for different numbers of bucket counters per bucket finding & scattering thread block, the thread block size is 256. For this benchmark we use the Geforce GTX480. Replicating the bucket counters results in significantly improved performance for non-uniformly distributed inputs.

## 6.4 GPU Sample Sort vs. revised GPU Sample Sort

On average, the revised implementation of GPU Sample Sort is approximately 25% faster than the original implementation on the Geforce GTX480. More important, it is up to 55% faster for input distributions on which our original implementation performs badly. See Section 5.5.2 for a description of the key distributions we use in our experiments.

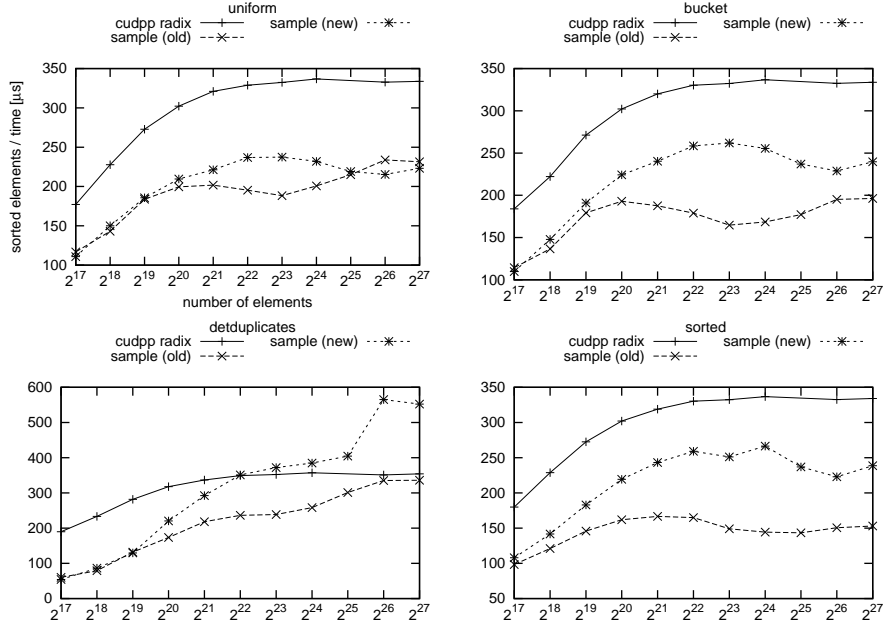


Figure 19: Original and revised GPU Sample Sorts performance on different 32 bit key distributions. As a reference we include performance numbers for CUDPPs radix sort and our merge sort. The better performance on sorted input is largely due to the additional bucket counter replication in the scattering kernel (step 4 of the  $k$ -way distribution algorithm).

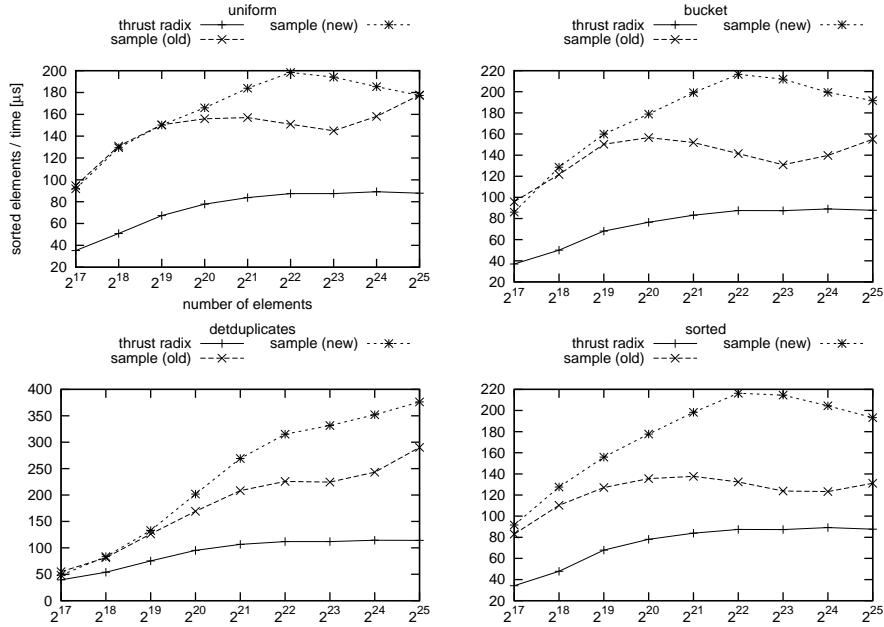


Figure 20: Original and revised GPU Sample Sorts performance on different 64 bit key distributions. Thrust radix sort is used as a reference; it is the only competitive GPU sorting implementation for 64 bit keys available to us.

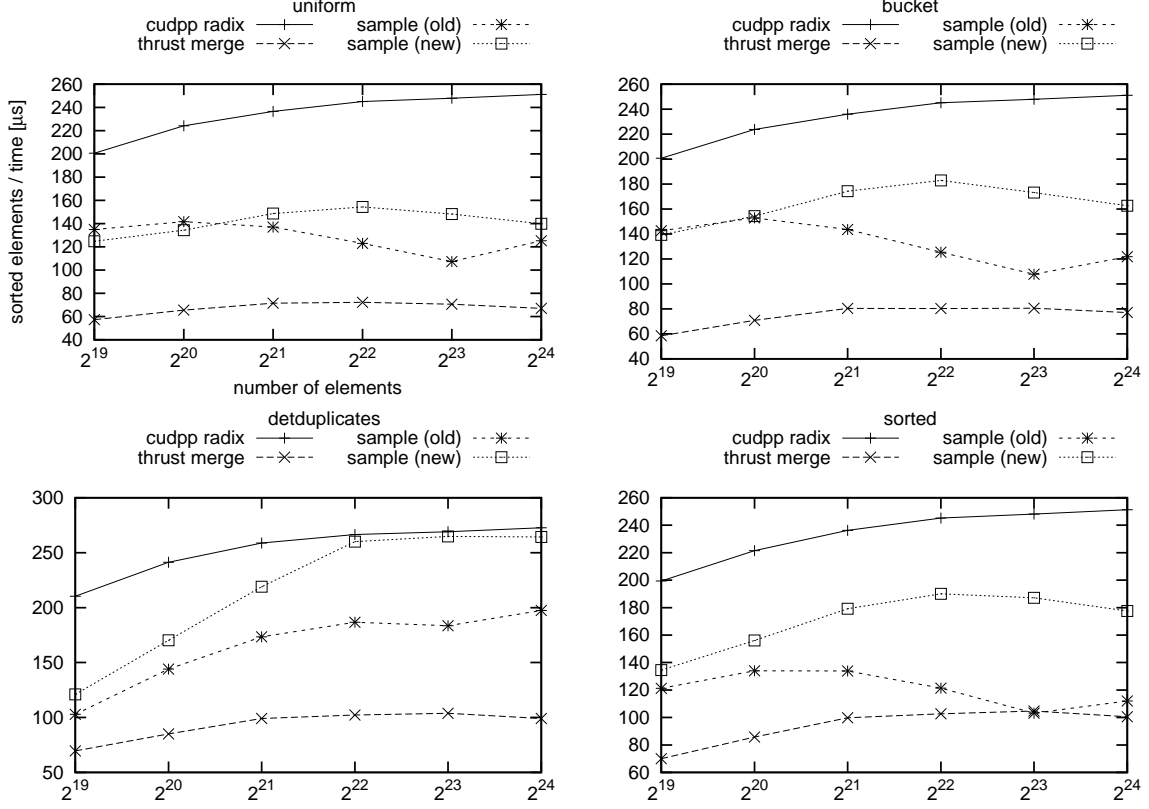


Figure 21: Original and revised GPU Sample Sorts performance on different 32 bit key distributions with 32 bit values.

## 6.5 Other possible implementation strategies

Our implementation of sample sort implicitly assumes that the parameter  $k$  is chosen as large as possible within memory limitations, to require as few  $k$ -way distribution passes as possible. As established, scattered writing and sharing bucket counters between threads are the price we pay for maximizing the parameter  $k$ . Limiting  $k$  to a smaller value than we use – for current GPU architectures a choice could be 16 – would allow for different trade-offs. Counters could be kept in registers for each thread, and it would be feasible to locally order keys by bucket id prior to scattering them to achieve sequential write access, or use local memory buffers for each bucket. Such an implementation could be similar to current GPU radix sort implementations, which usually process 4 bits per pass, and would be an interesting topic for future work.

## 7 Converting from Cuda to OpenCL

In this section we shortly compare the OpenCL [2] programming language to Cuda [1], and describe our experience of implementing GPU Sample Sort in OpenCL. OpenCL is available for different kinds of multi-core and many-core platforms. We, however, only look at it in the context of GPU programming.



## 7.1 The OpenCL programming language

OpenCL is very similar to Cuda. GPU kernels are launched by an application that runs on a CPU host, which also controls memory transfers between its memory and that of the GPU. In contrast to Cuda, with OpenCL the programmer is responsible for invoking the compilation of kernels at running time. The compiled kernels can be cached, however, to avoid unnecessary re-compilation. Therefore an application has to distribute the source code of its OpenCL kernels in some form, or the developers have to include compiled kernels for all OpenCL implementations on which they intend to run their application. An advantage of running time compilation is that it allows for dynamic specialization of kernel code via preprocessor directives and preprocessor defines, which can be handed to the compiler.

The syntax for kernels is almost identical, except for key word names. OpenCL forbids pointer traversal in some places; instead of modifying a pointer one has to use a base address pointer together with an offset. It is also impossible to declare global memory pointers that reside in local memory. A more notable difference lies in OpenCLs synchronization primitives. While Cuda only has one that synchronizes both local and global memory between threads of a thread block, OpenCL allows to explicitly synchronize either only local memory, only global memory, or both. Last, Cuda offers C++ style templates in kernel code. With OpenCL it is possible to emulate template parameters mostly via preprocessor defines. While this more error-prone and results in code that is less readable, it is also more powerful because the preprocessor defines can be chosen at running time in contrast to Cudas template parameters which have to be evaluated at the applications compile-time. Features like atomic operations, which are not supported on all GPUs, are exposed in OpenCL via extensions.

Even though OpenCL is based on the same two-level work decomposition as Cuda is (thread blocks and threads), it allows the programmer to leave out the specification of the thread block size. One can just specify the total number of threads and leave it to the OpenCL runtime to determine a suitable thread block size.

## 7.2 An OpenCL implementation of GPU Sample Sort

As expected, converting Cuda kernels to OpenCL is straightforward. The biggest complication is our usage of template parameters; but we are able to replace all parameters which are relevant for performance with preprocessor defines. For longer kernels the Nvidia OpenCL compiler is, however, less efficient at register allocation than the Cuda compiler is: an exact conversion of our thread block quicksort requires 22 registers per thread in OpenCL – the Cuda version only requires 9. Removing one specific code optimization – a special case for the base case sorting network, where we can omit a loop in the odd-even merge sort when the number of elements to sort is smaller than or equal to the number of threads – does lower the register usage in OpenCL to 14 registers per thread. Further experimentation with the code leads to no additional improvements. The quicksort kernel suffers from higher register consumption since it is mostly bound by memory latency, and with a higher register usage per thread fewer threads can run on a multiprocessor. As a result the access latencies cannot be hidden as much via SMT – see Section 3.1 for a more detailed explanation.

In a few places we are hindered by errors in the current Nvidia OpenCL implementation. Declaring variables of the type `boolean` in local memory leads to compilation errors. Replacing them with integer variables, which should be semantically identical in the C99-based OpenCL, causes our code to malfunction. As a workaround we replace boolean variables with boolean *arrays* of size 1. Working with OpenCL is less comfortable than working with Cuda is, at the time of writing; the Nvidia OpenCL compiler may crash without giving any meaningful error messages in response to syntactic or semantic errors in kernel code.

The current Nvidia drivers with support for Fermi-based GPUs (e.g. the Geforce GTX480) lack stable OpenCL support. We are only able to compare an OpenCL implementation to a Cuda implementation on a different system, which has an older driver and a different GPU. This test system has an Intel E8500 CPU with 8 GiB RAM and a Gainward Geforce 9800GT 1024 MiB. It runs Windows 7 64 bit and the ForceWare 197.45 GPU driver. The Cuda 3.0 compiler is used together with the Visual Studio 2008 C++ compiler. Input data is generated in the same way as described in Section 5.5.2.

Since the Geforce 9800GT lacks atomic operations, we use a different variant of our GPU Sample Sort implementation. The main difference is that its second level sorter is quicksort, like in our originally published version, and that it serializes the bucket counter incrementing in the bucket finding and scattering kernel – refer to Section 6.1 and Section 6.2 for a description of GPU Sample Sorts kernels. The OpenCL and Cuda implementations use identical parameters, except for quicksort, where we use 128 threads per block instead of 256 to counter the negative effect of the increased register consumption of the OpenCL version. Since we use the scan and reduction primitives from the Thrust [34] library in our Cuda implementation, we also have to implement replacements for them in OpenCL. To ensure comparability between the OpenCL and Cuda implementation we also implement our own scan and reduction primitive in Cuda and use these instead of the ones from Thrust.

On average, the OpenCL implementation reaches 88.7% of the Cuda implementations performance, see Figure 22. Examining the GPU running times shows a smaller difference between the OpenCL and Cuda kernels; we conclude that the CPU overhead for launching kernels is currently larger for OpenCL. Copying data between CPU memory and GPU memory is also 5% to 15% slower with OpenCL. We expect that – as the OpenCL compiler and runtime implementation becomes more mature – the difference in performance between OpenCL and Cuda will decrease further.

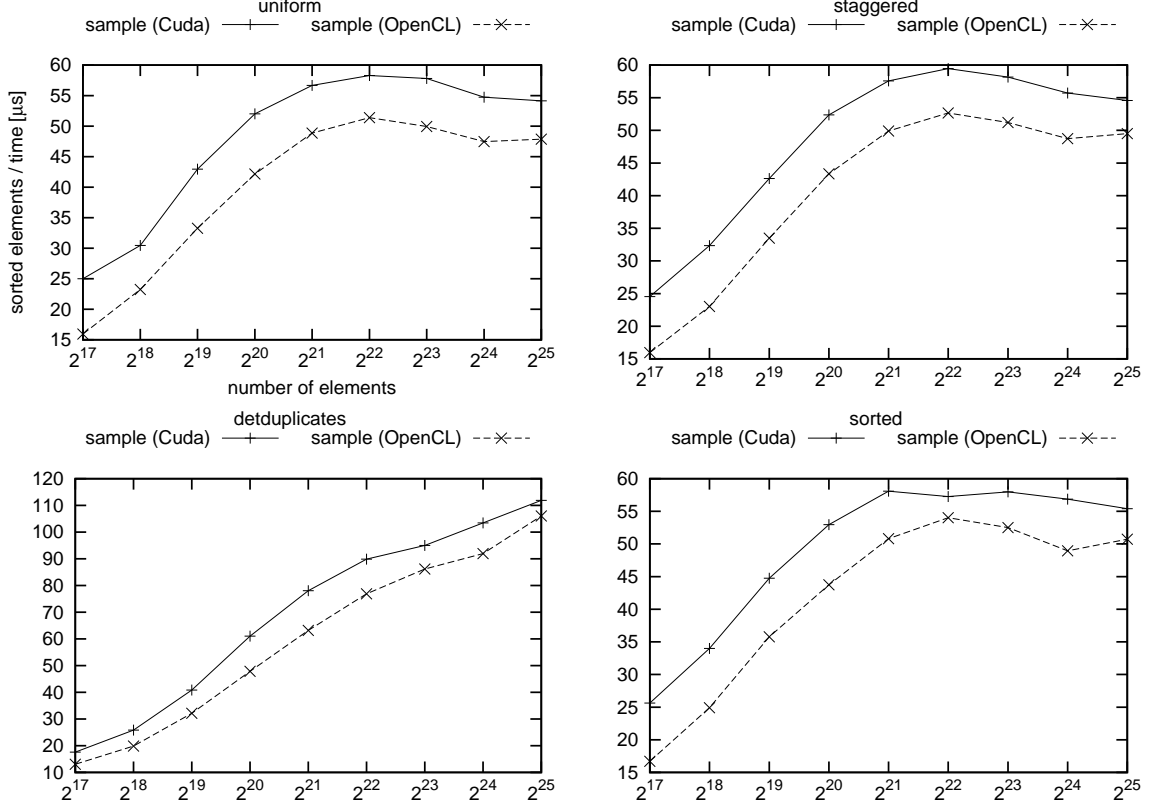


Figure 22: Cuda and OpenCL GPU Sample Sorts performance on different 32 bit key distributions. These benchmarks are performed with a Geforce 9800GT and a modified version of GPU Sample Sort that does not use atomic operations, consequently the sorting rates are much lower than in our other experiments. The OpenCL implementation behaves the same as the Cuda implementation but generally performs worse.

## 8 Conclusion

Our work on merging resulted in a merging primitive with asymptotically optimal work complexity and a running time faster than or similar to other GPU merging routines, except for a hypothetical merging primitive derived from the merge sort algorithm by Satish et al. [35]. We also demonstrated that sorting and merging are computationally bound on current generation GPUs. Our multi sequence selection algorithms show that step complexity should not be over-rated when analyzing GPU algorithms. The selection algorithm due to Vitaly Osipov also is an example for an algorithm which can be practical on the GPU even though it contains a substantial amount of sequential work. Further, our investigation of merging on GPUs made us discover and re-discover several open problems.

As we noted in Section 5.3.2, the search for practical merging networks with a better work complexity than the bitonic and the odd-even network has not yielded results for decades. The discovery of the AKS network [3] with a work complexity of  $\mathcal{O}(n \cdot \log(n))$  suggests that better networks for merging and sorting may exist. However, the only ad-

vances that have been made are networks which lend themselves to more flexible hardware implementations – for example the balanced sorting network [14].

Multi-way merging on GPUs also warrants closer investigation. The four-way merging routine we presented in Section 5.3.5 is certainly not the only feasible implementation strategy.

In Section 6 we demonstrated incremental improvements to GPU Sample Sort that lead to 25% higher average performance and elaborated on possible different implementation strategies for sample sort.

We were successful at porting our implementation to OpenCL, achieving 88.7% of our Cuda implementations performance. We concluded that OpenCL can be a replacement for Cuda, although it is not yet as mature and stable as Cuda is. We only tested our OpenCL implementation on Nvidia hardware; it would be interesting to investigate the performance of GPU Sample Sort on other platforms that support OpenCL.

Lastly, we argued that there is a lot of work to be done on the topic of abstract machine models and programming strategies for many-core processors. Finding design strategies which lead to algorithms that can be adapted and tuned easily for different kinds of architectures is an important goal. Thus far Merril and Grimshaw [27, 28] described the most rigorous approach to the tuning of GPU algorithms.

# A Proof of correctness for sequential multi sequence selection

## A.1 Correctness

To establish the correctness of Algorithm 1, we show that  $U$  and  $L$ , which always represent the upper and lower boundary for the size of the current selection  $c$ , converge to the target size  $t \in [0, n]$ , where  $n$  is the combined size of all sequences. The precondition for our proof is that the input sequences are *sorted*, and that in each step a valid sequence for increasing or decreasing the current selection is chosen. We show that after each step after the first step of the algorithm either  $L \leq t < U$ ,  $L < U \leq t$  or  $L = U = t$  holds true.




Figure 23: In this case the selection is increased and the previous position in the sequence is added to the lower boundary.

Let  $L < U \leq t$ ; we pick a sequence to *increase* the selection with, see Figure 23.  $L$  is increased by the difference between the last position in that sequence and the lower bound of the sequence.  $U$  is increased by the difference between the new position and the last position. Thus,  $L$  is now at most as large as  $U$  was previously.  $L$  can only become equal to  $U$  if  $U$  is not increased in the step. Also, for  $L$  to become equal to  $U$ ,  $U$  must be equal to  $t$ . If it is not, then we have either not chosen a valid sequence or  $t$  is larger than the total number of elements  $n$ . If  $L$  is still smaller than  $U$  then either  $L \leq t < U$  or  $L < U \leq t$  hold true, because  $U \leq t$  held true prior to the step.




Figure 24: We decrease the size of the selection. The global upper boundary of the sequence is decreased.

Let  $L \leq t < U$ ; we pick a sequence to *decrease* the selection with, see Figure 24.  $L$  is increased by the difference between the position of the greatest element in the sequence that belongs to the final selection. Thus,  $L \leq t$  still holds true and therefore also either  $L \leq t < U$ ,  $L < U \leq t$  or  $L = U = t$ .

If  $L = U = t$  is true the algorithm terminates. The local upper boundaries  $u_1 \dots u_k$  decrease monotonically.  $U$  cannot grow smaller than  $L$  or larger than the sum of the local upper boundaries. The local lower boundaries  $l_1 \dots l_k$  increase monotonically.  $L$  cannot grow larger than  $t$  or smaller than the sum of the local lower boundaries. It follows that  $U$  will converge to  $L$ , the algorithm will terminate after a finite number of steps, and it will terminate with a selection of size  $t$ .

## A.2 Consistency and stability

As established in Section 5.2, selections may be ambiguous in the presence of duplicate keys. To be exact, multiple selections of a given size exist if the greatest key belonging to a selection of this size occurs in more than one sequence. We achieve consistent and stable selections by breaking ties – choosing a sequence when there are multiple sequences whose sample element has the same key and that key is the smallest or greatest of all sample element keys – by choosing the sequence with the lowest index for minimum selection, and choosing the sequence with the greatest index for maximum selection.




Figure 25: A tie between two sequences;  $\text{key}_{\max}$  is the greatest key that is in the selection and it occurs in two sequences. The bars represent the current selections in these sequences. All instances of  $\text{key}_{\max}$  in  $\text{sequence}_i$  are selected before any in  $\text{sequence}_j$  are selected. Vice versa, all instances in  $\text{sequence}_j$  are deselected before any in  $\text{sequence}_i$  are deselected.


Assume  $\text{key}_{\max}$  is the greatest key that is in the selection and  $m$  instances of  $\text{key}_{\max}$  are in the selection. The sequences are numbered by their order of appearance in the input. We look at the case where  $\text{key}_{\max}$  occurs  $x$  times in  $\text{sequence}_i$ ,  $y$  times in  $\text{sequence}_j$ , where  $i < j$ , and in no other sequence. It follows from the correctness of the algorithm that eventually all other sequences are settled and only the boundaries for  $\text{sequence}_i$  and  $\text{sequence}_j$  remain to be determined. If  $m \leq x$  and any instance of  $\text{key}_{\max}$  in  $\text{sequence}_j$  is selected, then all instances of  $\text{key}_{\max}$  are selected in  $\text{sequence}_i$  because  $\text{sequence}_i$  takes precedence over  $\text{sequence}_j$  for minimum selection and thus it takes precedence for increasing the selection size. It follows that more than  $m$  instances of  $\text{key}_{\max}$  are selected and the selection size is therefore decreased. Here  $\text{sequence}_j$  takes precedence over  $\text{sequence}_i$ ; all instances of  $\text{key}_{\max}$  in  $\text{sequence}_j$  are deselected before any in  $\text{sequence}_i$  are. Thus, if  $m \leq x$ , all instances of  $\text{key}_{\max}$  that are selected when the algorithm terminates lie in  $\text{sequence}_i$ . If  $m > x$ , then all instances of  $\text{key}_{\max}$  in  $\text{sequence}_i$  are selected and some in  $\text{sequence}_j$ . Because  $\text{sequence}_i$  precedes  $\text{sequence}_j$  the selection can be used for stable merging. Also, the selection is a superset of any smaller selection and is therefore consistent. The same reasoning applies for ties between more than two sequences; they are

resolved by successively resolving ties between pairs of sequences in the manner described here.

## B Proof of correctness for block-wise merging

Let  $X$  and  $Y$  denote the sorted input sequences,  $x_i$  the element from  $X$  at position  $i$  and  $y_j$  the element from  $Y$  at position  $j$ . We load blocks of size  $B$  into a buffer of size  $2 \cdot B$ . Let  $x_a$  denote the element behind the most recent block from  $X$ , and  $y_b$  the next element from  $Y$ . After each merge step we output the lower half of the buffer and load a block from  $X$  if  $x_a \leq y_b$ . Otherwise we load the next block from  $Y$ . To show that the procedure merges  $X$  and  $Y$  correctly we have to show that the buffer always contains the next smallest  $B$  elements. We prove this by induction.

Case 1:



Case 2:




Figure 26: In the first case both initial blocks are taken from the same sequence, in the second one is taken from each.

Without loss of generality we differentiate between two base cases which are illustrated in Figure 26. In the first, the upper half of the buffer is initialized with a block from  $X$  and the next block is also taken from  $X$ . Thus,  $x_B \leq y_0$  and therefore the buffer contains the smallest  $B$  elements. In the second case the first block is taken from  $X$  and the second one from  $Y$ . It follows that the buffer contains  $B$  elements from each sequence. We need at most  $B$  elements from  $X$  – iff  $x_{B-1} \leq y_0$  – and vice versa we need at most  $B$  elements from  $Y$ . Since the buffer contains exactly  $B$  elements from each sequence it always holds the smallest  $B$  elements.




Figure 27: The current position in  $X$  and  $Y$  is denoted by  $x_a$  and  $y_b$ . Two blocks of size  $B$  fit into the buffer. The upper half of the buffer contains the elements  $x_{a-u} \dots x_{a-1}$  and  $y_{b-v} \dots y_{b-1}$ , where  $u + v = B$ . A new block is taken from  $X$ , since  $x_a \leq y_b$ .

For the inductive step we assume that the merge buffer up to this point always contained the next  $B$  smallest elements. Therefore, for both sequences the distance between the previous  $B$  smallest elements and the current position in that sequence can be at most  $B$ . Otherwise, the previous merge step would have been incorrect. Let  $u$  denote that distance for  $X$  and  $v$  for  $Y$ . Trivially,  $u + v = B$ . Without loss of generality, assume that the next block is taken from  $X$ . The next merge will yield an incorrect result only if more than  $v$  elements from  $Y$  are among the next  $B$  smallest elements. If more than  $v$  elements from  $Y$  were needed, then  $x_a > y_b$ . Thus, we would not have taken the next block from  $X$ . This is illustrated in Figure 27. By induction it follows that the buffer will always contain the next  $B$  smallest elements. Thus we have shown the correctness of block-wise merging.

## C CREW PRAM Quicksort

Algorithm 7 is based on the GPU quicksort by Cederman and Tsigas [9]. It operates on global memory and switches to a different sorting strategy when a partition is small enough to fit into local multiprocessor memory.



---

**Algorithm 7:** Parallel Quicksort.

---

```
/* Let  $i \in 0 \dots p-1$  and  $p$  denote the local thread index and thread
   count. */
QSort( $e = \langle e_1, \dots, e_n \rangle$ )
begin
  if  $n < M$  then return SmallSort( $e$ )
  if  $i = 0$  then pick a pivot  $d \in e$ 
   $a_i := b_i := c_i := 0$ 
   $t := e$ 
  for  $i \frac{n}{p} \leq j \leq (i+1) \frac{n}{p}$  do
    if  $t_j < d$  then  $a_i++$ 
    else if  $t_j = d$  then  $b_i++$ 
    else  $c_i++$ 
  end
  scan( $a_0 \dots a_{p-1}$ )
  scan( $b_0 \dots b_{p-1}$ )
  scan( $c_0 \dots c_{p-1}$ )
  for  $i \frac{n}{p} \leq j \leq (i+1) \frac{n}{p}$  do
    if  $t_j < d$  then  $e[a_i] := t_j$   $a_i++$ 
    else if  $t_j = d$  then  $e[b_i] := t_j$   $b_i++$ 
    else  $e[c_i] := t_j$   $c_i++$ 
  end
  return
  Concatenate(QSort( $\langle e_1, \dots, e_{a_i} \rangle$ ),  $\langle e_{a_i+1}, \dots, e_{b_i} \rangle$ , QSort( $\langle e_{b_i+1}, \dots, e_n \rangle$ ))
end
```

---

## References

- [1] CUDA. <http://www.nvidia.com/CUDA/>.
- [2] OpenCL. <http://www.khronos.org/opencv/>.
- [3] M. Ajtai, J. Komlós, and E. Szemerédi. An  $O(n \log n)$  sorting network. In *STOC '83: Proceedings of the fifteenth annual ACM symposium on Theory of computing*, pages 1–9, New York, NY, USA, 1983. ACM.
- [4] R. Baraglia, G. Capannini, F. M. Nardini, and F. Silvestri. Sorting using Bitonic network with CUDA. Universidade Tecnica de Lisboa, 2009.
- [5] K. E. Batcher. Sorting networks and their applications. In *AFIPS '68 (Spring): Proceedings of the April 30–May 2, 1968, spring joint computer conference*, pages 307–314, New York, NY, USA, 1968. ACM.
- [6] G. E. Blelloch. Scans as primitive parallel operations. *IEEE Trans. Comput.*, 38(11):1526–1538, 1989.
- [7] G. E. Blelloch. Prefix sums and their applications. Technical Report CMU-CS-90-190, School of Computer Science, Carnegie Mellon University, Nov. 1990.
- [8] G. E. Blelloch, C. E. Leiserson, B. M. Maggs, C. G. Plaxton, S. J. Smith, and M. Zagha. A comparison of sorting algorithms for the connection machine CM-2. In *SPAA '91: Proceedings of the third annual ACM symposium on Parallel algorithms and architectures*, pages 3–16, New York, NY, USA, 1991. ACM.
- [9] D. Cederman and P. Tsigas. GPU-Quicksort: A practical Quicksort algorithm for graphics processors. *J. Exp. Algorithmics*, 14:1.4–1.24, 2009.
- [10] J. Chhugani, A. D. Nguyen, V. W. Lee, W. Macy, M. Hagog, Y.-K. Chen, A. Baransi, S. Kumar, and P. Dubey. Efficient implementation of sorting on multi-core SIMD CPU architecture. *PVLDB*, 1(2):1313–1324, 2008.
- [11] R. Cole. Parallel merge sort. *SIAM J. Comput.*, 17(4):770–785, 1988.
- [12] F. Dehne and H. Zaboli. Deterministic Sample Sort For GPUs. *CoRR*, abs/1002.4464, 2010.
- [13] S. Ding, J. He, H. Yan, and T. Suel. Using graphics processors for high-performance IR query processing. In *WWW '08: Proceeding of the 17th international conference on World Wide Web*, pages 1213–1214, New York, NY, USA, 2008. ACM.
- [14] M. Dowd, Y. Perl, M. Saks, and L. Rudolph. The balanced sorting network. In *PODC '83: Proceedings of the second annual ACM symposium on Principles of distributed computing*, pages 161–172, New York, NY, USA, 1983. ACM.
- [15] R. S. Francis, I. D. Mathieson, and L. Pannan. A fast, simple algorithm to balance a parallel multiway merge. In *PARLE '93: Proceedings of the 5th International PARLE Conference on Parallel Architectures and Languages Europe*, pages 570–581, London, UK, 1993. Springer-Verlag.

- [16] M. Frigo, C. E. Leiserson, H. Prokop, and S. Ramachandran. Cache-oblivious algorithms. In *FOCS '99: Proceedings of the 40th Annual Symposium on Foundations of Computer Science*, Washington, DC, USA, 1999. IEEE Computer Society.
- [17] L. Ha, J. Krüger, and C. Silva. Fast Four-Way Parallel Radix Sorting on GPUs. In *Computer Graphics Forum*, volume 28, pages 2368–2378. Blackwell Publishing, Dec. 2009.
- [18] T. Hagerup and C. Rüb. Optimal merging and sorting on the EREW PRAM. *Inf. Process. Lett.*, 33(4):181–185, 1989.
- [19] D. R. Helman, D. A. Bader, and J. JáJá. A randomized parallel sorting algorithm with an experimental study. *J. Parallel Distrib. Comput.*, 52(1):1–23, 1998.
- [20] B. Huang, J. Gao, and X. Li. An Empirically Optimized Radix Sort for GPU. *Parallel and Distributed Processing with Applications, International Symposium on*, 0:234–241, 2009.
- [21] H. Inoue, T. Moriyama, H. Komatsu, and T. Nakatani. AA-Sort: A New Parallel Sorting Algorithm for Multi-Core SIMD Processors. In *PACT '07: Proceedings of the 16th International Conference on Parallel Architecture and Compilation Techniques*, pages 189–198, Washington, DC, USA, 2007. IEEE Computer Society.
- [22] M. F. Ionescu. Optimizing parallel bitonic sort. In *IPPS '97: Proceedings of the 11th International Symposium on Parallel Processing*, pages 303–309, Washington, DC, USA, 1997. IEEE Computer Society.
- [23] D. E. Knuth. *The art of computer programming, volume 3: (2nd ed.) sorting and searching*. Addison Wesley Longman Publishing Co., Inc., Redwood City, CA, USA, 1998.
- [24] H. W. Lang. Sorting networks, 2008.  
[www.inf.fh-flensburg.de/lang/algorithmen/sortieren/](http://www.inf.fh-flensburg.de/lang/algorithmen/sortieren/).
- [25] N. Leischner, V. Osipov, and P. Sanders. GPU Sample Sort. In *Proc. of the 24th International Symposium on Parallel and Distributed Processing (IPDPS)*. IEEE Computer Society, 2010.
- [26] M. Matsumoto and T. Nishimura. Mersenne twister: A 623-dimensionally equidistributed uniform pseudo-random number generator. *ACM Transactions on Modeling and Computer Simulation*, 8(1):3–30, 1998.
- [27] D. Merrill and A. Grimshaw. Parallel scan for stream architectures. Technical report, University of Virginia, Department of Computer Science, Charlottesville, VA, USA, 2009.
- [28] D. Merrill and A. Grimshaw. Revisiting Sorting for GPGPU Stream Architectures. Technical Report CS2010-03, University of Virginia, Department of Computer Science, Charlottesville, VA, USA, 2010.
- [29] D. R. Musser. Introspective sorting and selection algorithms. *Softw. Pract. Exper.*, 27(8):983–993, 1997.
- [30] L. Natvig. Logarithmic time cost optimal parallel sorting is not yet fast in practice! In *Supercomputing '90: Proceedings of the 1990 ACM/IEEE conference on Supercomputing*, pages 486–494, Los Alamitos, CA, USA, 1990. IEEE Computer Society Press.

- [31] H. Peters, O. Schulz-Hildebrandt, and N. Luttenberger. Fast comparison-based in-place sorting with CUDA. In *Eighth International Conference on Parallel Processing and Applied Mathematics*, Sept. 2009.
- [32] F. Putze, P. Sanders, and J. Singler. MCSTL: the multi-core standard template library. In *PPoPP '07: Proceedings of the 12th ACM SIGPLAN symposium on Principles and practice of parallel programming*, pages 144–145, New York, NY, USA, 2007. ACM.
- [33] P. Sanders and S. Winkel. Super scalar sample sort. In *Proc. of the 12th European Symposium on Algorithms (ESA)*, volume 3221 of *LNCS*, pages 784–796. Springer, 2004.
- [34] N. Satish, M. Harris, and M. Garland. Designing efficient sorting algorithms for manycore GPUs. In *IPDPS '09: Proceedings of the 23th IEEE International Symposium on Parallel&Distributed Processing*, pages 1–10, Washington, DC, USA, 2009. IEEE Computer Society.
- [35] N. Satish, C. Kim, J. Chhugani, A. D. Nguyen, V. W. Lee, D. Kim, and P. Dubey. Fast sort on CPUs and GPUs: a case for bandwidth oblivious SIMD sort. Technical report, New York, NY, USA, 2010.
- [36] S. Sengupta, M. Harris, Y. Zhang, and J. D. Owens. Scan primitives for GPU computing. In *Proc. ACM SIGGRAPH/EUROGRAPHICS Conference on Graphics Hardware*, pages 97–106, 2007.
- [37] E. Sintorn and U. Assarsson. Fast parallel GPU-sorting using a hybrid algorithm. *J. Parallel Distrib. Comput.*, 68(10):1381–1388, 2008.
- [38] P. J. Varman, S. D. Scheufler, B. R. Iyer, and G. R. Ricard. Merging multiple lists on hierarchical-memory multiprocessors. *J. Parallel Distrib. Comput.*, 12(2):171–177, 1991.
- [39] H. Wong, M.-M. Papadopoulou, M. Sadooghi-Alvandi, and A. Moshovos. Demystifying GPU Microarchitecture through Microbenchmarking. In *In Proc. IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS), March 2010*, Washington, DC, USA, 2010. IEEE Computer Society.
- [40] X. Ye, D. Fan, W. Lin, N. Yuan, and P. Ienne. High Performance Comparison-Based Sorting Algorithm on Many-Core GPUs. In *Proc. of the 24th International Symposium on Parallel and Distributed Processing (IPDPS)*. IEEE Computer Society, 2010.